



LFS201

Essentials of System Administration

Version 1.1



LFS201: Version 1.0

© Copyright the Linux Foundation 2015. All rights reserved.

© Copyright the Linux Foundation 2015. All rights reserved.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the Linux Foundation
<http://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact:

training@linuxfoundation.org

Contents

1	Preface	1
2	System Startup and Shutdown	3
3	GRUB	5
4	init: SystemV, Upstart, systemd	7
5	Linux Filesystem Tree Layout	11
6	Kernel Services and Configuration	15
7	Kernel Modules	17
8	Devices and udev	19
9	Partitioning and Formatting Disks	21
10	Encrypting Disks	27
11	Linux Filesystems and the VFS	31
12	Filesystem Features: Attributes, Creating, Checking, Mounting	33
13	Filesystem Features: Swap, Quotas, Usage	37
14	The Ext2/Ext3/Ext4 Filesystems	41
15	The XFS and btrfs Filesystems	45
16	Logical Volume Management (LVM)	47
17	RAID	49
18	Local System Security	51
19	Linux Security Modules	55
20	Processes	59
21	Signals	63

22 System Monitoring	67
23 Process Monitoring	69
24 I/O Monitoring and Tuning	71
25 I/O Scheduling	75
26 Memory: Monitoring Usage and Tuning	79
27 Package Management Systems	81
28 RPM	83
29 DPKG	87
30 yum	89
31 zypper	93
32 APT	95
33 User Account Management	99
34 Group Management	103
35 File Permissions and Ownership	105
36 Pluggable Authentication Modules (PAM)	107
37 Backup and Recovery Methods	109
38 Network Addresses	113
39 Network Devices and Configuration	115
40 Firewalls	119
41 Basic Troubleshooting	123
42 System Rescue	125

Chapter 1

Preface



Lab 1.1: Configuring the System for `sudo`

It is very dangerous to run a **root shell** unless absolutely necessary: a single typo or other mistake can cause serious (even fatal) damage.

This lab describes the procedure for configuring things such that simple commands may be run with superuser privileges by using the **sudo** mechanism. With **sudo** the user only needs to know their own password and never needs to know the root password.

If you are using a distribution such as **Ubuntu**, you may not need to do this lab to get **sudo** configured properly for the course. However, you should still make sure you understand the procedure.

To check if your system is already configured to let the user account you are using run **sudo**, just do a simple command like:

```
$ sudo ls
```

You should be prompted for your user password and then the command should execute. If instead, you get an error message you need to execute the following procedure.

Launch a root shell by typing **su** and then giving the **root** password, not your user password.

On all recent **Linux** distributions you should navigate to the `/etc/sudoers.d` subdirectory and create a file, usually with the name of the user to whom root wishes to grant **sudo** access. However, this convention is not actually necessary as **sudo** will scan all files in this directory as needed. The file can simply contain:

```
student ALL=(ALL) ALL
```

if the user is **student**.

An older practice (which certainly still works) is to add such a line at the end of the file `/etc/sudoers`. It is best to do so using the **visudo** program, which is careful about making sure you use the right syntax in your edit.

You probably also need to set proper permissions on the file by typing:

```
$ chmod 440 /etc/sudoers.d/student
```

(Note some **Linux** distributions may require 400 instead of 440 for the permissions.)

After you have done these steps, exit the root shell by typing `exit` and then try to do `sudo ls` again.

There are many other ways an administrator can configure **sudo**, including specifying only certain permissions for certain users, limiting searched paths etc. The `/etc/sudoers` file is very well self-documented.

However, there is one more setting we highly recommend you do, even if your system already has **sudo** configured. Most distributions establish a different path for finding executables for normal users as compared to root users. In particular the directories `/sbin` and `/usr/sbin` are not searched, since **sudo** inherits the `PATH` of the user, not the full root user.

Thus, in this course we would have to be constantly reminding you of the full path to many system administration utilities; any enhancement to security is probably not worth the extra typing and figuring out which directories these programs are in. Consequently, we suggest you add the following line to the `.bashrc` file in your home directory:

```
PATH=$PATH:/usr/sbin:/sbin
```

If you log out and then log in again (you don't have to reboot) this will be fully effective.

Chapter 2

System Startup and Shutdown



Lab 2.1: Shutdown VS. Halt VS. Reboot

NOTE: This exercise requires that it be run from the console (i.e., not over the network through SSH).

1. Reboot the system using **shutdown**.
2. Power off the system using **shutdown**.
3. Power the system back up.

Solution 2.1

1. `$ sudo shutdown -r now`
2. `$ sudo shutdown -h now`
3. Press the power button , or restart your virtual machine.

Chapter 3

GRUB



Lab 3.1: Booting into Non-Graphical Mode Using GRUB

NOTE: This exercise requires that it be run from the console (i.e., not over **SSH**).

1. Reboot your machine and go into the **GRUB** interactive shell by hitting **e** (or whatever other key is required as listed on your screen.)
2. Make your system boot into non-graphical mode. How you do this depends on the system.
On traditional systems that respect **runlevels** (which we will talk about in the next section) you can append a **3** to the kernel command line in the specific entry you pick from the **GRUB** menu of choices.
On some other systems (including **Ubuntu**) you need to append **text** instead.
3. Hit the proper key to make system continue booting.
4. After the system is fully operational in non-graphical mode, bring it up to graphical mode. Depending on your system, one of the following commands should do it:

```
$ sudo telinit 5
$ sudo service gdm restart
$ sudo service lightdm restart
```


Chapter 4

init: SystemV, Upstart, systemd



Lab 4.1: Adding a New Startup Service with SysVinit

In this and the following exercise, we will create a simple startup service. First we will do it for a **SysVinit** system. Note that if you are using a **systemd**-based system everything should still work because of the backwards compatibility layer that all distributions utilize. However, in the next exercise we will do natively for **systemd**.

If you are on a **Debian**-based system like **Ubuntu**, make sure you have installed the **sysvinit-utils** and **chkconfig** packages. However, recent versions of **Ubuntu** no longer package **chkconfig**; you'll have to use the **update-rc.d** utility instead.

First we have to create the service-specific script; you can create one of your own for fun, or to get the procedure down just (as root) create a file named `/etc/init.d/fake_service` containing the following content:

```
#!/bin/bash
# fake_service
# Starts up, writes to a dummy file, and exits
#
# chkconfig: 35 69 31
# description: This service doesn't do anything.
# Source function library

. /etc/sysconfig/fake_service

case "$1" in
start) echo "Running fake_service in start mode..."
    touch /var/lock/subsys/fake_service
    echo "$0 start at $(date)" >> /var/log/fake_service.log
    if [ ${VAR1} = "true" ]
    then
        echo "VAR1 set to true" >> /var/log/fake_service.log
    fi
    echo
;;
```

```

stop)
echo "Running the fake_service script in stop mode..."
echo "$0 stop at $(date)" >> /var/log/fake_service.log
if [ ${VAR2} = "true" ]
then
echo "VAR2 = true" >> /var/log/fake_service.log
fi
rm -f /var/lock/subsys/fake_service
echo
;;
*)
echo "Usage: fake_service {start | stop}"
exit 1
esac
exit 0

```

If you are taking the online self-paced version of this course, the script is available for download from your **Lab** screen. Make the file above executable and give other proper permissions:

```
$ sudo chmod 755 /etc/init.d/fake_service
```

You'll notice the script includes the file `/etc/sysconfig/fake-service`. (On non-**RHEL** systems you should change this to `/etc/default/fake_service`.) Create it and give it the following contents:

```

VAR1="true"
VAR2="true"

```

Test to see if the script works properly by running the following commands:

```

$ sudo service fake_service
$ sudo service fake_service start
$ sudo service fake_service stop

```

Look at the file named `/var/log/fake-service.log`. What does it contain?

For fun you can add additional modes like `restart` to the script file; look at other scripts in the directory to get examples of what to do.

Next we will want to have the ability to start **fake-service** whenever the system starts, and stop it when it shuts down. If you do:

```
$ sudo chkconfig --list fake_service
```

you will get an error as it hasn't been set up yet for this. You can easily do this with:

```
$ sudo chkconfig --add fake_service
```

and you can turn it on or off at boot time with

```

$ sudo chkconfig fake_service on
$ sudo chkconfig fake_service off

```

To test this completely you'll have to reboot the system to see if it comes on automatically. You can also try varying the runlevels in which the service is running.

Lab 4.2: Adding a New Startup Service with systemd

As mentioned in the previous exercise, you can still use the **SysVinit** startup script procedure with **systemd** but this is deprecated.

The analogous procedure is to create (as root) a file directly under `/etc/systemd/system` or somewhere else in that directory tree; distributions have some varying tastes on this. For example a very minimal file named `/etc/systemd/system/fake2.service`

```
[Unit]
Description=fake2
After=network.target

[Service]
ExecStart=/bin/echo I am starting the fake2 service
ExecStop=/bin/echo I am stopping the fake2 service

[Install]
WantedBy=multi-user.target
```

Now there are many things that can go in this **unit** file. The `After=network.target` means the service should start only after the network does, while the `WantedBy=multi-user.target` means it should start when we reach multiple-user mode. This is equivalent to runlevels 2 and 3 in **SysVinit**. Note `graphical.target` would correlate with runlevel 5.

Change the permissions on the file to make it executable:

```
$ chmod 755 /etc/systemd/system/fake2.service
```

Now all we have to do to start, stop and check the service status are to issue the commands:

```
$ sudo systemctl start fake2.service
$ sudo systemctl status fake2.service
$ sudo systemctl stop fake2.service
```

If you are fiddling with the unit file while doing this you'll need to reload things with:

```
$ sudo systemctl daemon-reload
```

as the system will warn you.

To set things up so the service turns on or off on system boot:

```
$ sudo systemctl enable fake2.service
$ sudo systemctl disable fake2.service
```

Once again, you really need to reboot to make sure it has taken effect.

Chapter 5

Linux Filesystem Tree Layout



Lab 5.1: Sizes of the Default Linux Directories

Use the **du** utility to calculate the overall size of each of your system's top-level directories.

Type the command:

```
$ du --help
```

for hints on how to obtain and display this result efficiently.

Solution 5.1

To obtain a full list of directories under **/** along with their size:

```
$ sudo du --max-depth=1 -hx /
```

```
4.3M    /home
16K     /lost+found
39M     /etc
4.0K    /srv
3.6M    /root
178M    /opt
138M    /boot
6.1G    /usr
1.1G    /var
16K     /mnt
4.0K    /media
869M    /tmp
8.4G    /
```

Where we have used the options:

- `--maxdepth=1`: Just go down one level from `/` and sum up everything recursively underneath in the tree.
- `-h`: Give human-readable numbers (KB, MB, GB).
- `-x` Stay on one filesystem; don't look at directories that are not on the `/` partition. In this case that means ignore:

`/dev /proc /run /sys`

because these are pseudo-filesystems which exist in memory only; they are just empty mount points when the system is not running. Because this is a **RHEL 7** system, the following mount points are also not followed:

`/bin /sbin /lib /lib64`

since they are just symbolically linked to their counterparts under `/usr`.

Lab 5.2: Touring the `/proc` Filesystem

Exactly what you see in this exercise will depend on your kernel version, so you may not match the output shown precisely.

1. As root, `cd` into `/proc` and do a directory listing. This should display a number of files and directories:

```
$ cd /proc
$ ls -F
1/      17/    2180/   2541/  34/     508/    636/   773/           locks
10/     1706/   22/     259/   3469/   510/    644/   794/           meminfo
1009/   1707/   2203/   26/    35/     512/    645/   8/            misc
1014/   1775/   2231/   2626/  36/     513/    66/    825/           modules
1015/   1779/   2233/   263/   37/     515/    67/    826/           mounts@
1019/   18/     2234/   2635/  374/    517/    676/   879/           mtrr
1023/   1846/   2241/   264/   3792/   519/    68/    9/            net@
11/     1898/   23/     266/   3857/   521/    681/   ACPI/         pagetypeinfo
1144/   19/     2319/   27/    3858/   5217/   6824/   asound/        partitions
12/     1901/   2323/   271/   3865/   537/    6909/   buddyinfo     sched_debug
1242/   1905/   2337/   278/   3866/   538/    6979/   bus/           schedstat
1265/   1908/   2338/   279/   395/    555/    7/      cgroups        scsi/
1295/   1923/   2363/   28/    397/    556/    7053/   cmdline        self@
1296/   1931/   238/    2897/  3990/   5564/   7091/   config.gz      slabinfo
1297/   1935/   239/    29/    409/    5571/   7123/   consoles       softirqs
1298/   1941/   23957/  2928/  42/     5768/   7188/   cpuinfo        stat
1299/   2/      24/     2945/  43/     583/    7222/   crypto         swaps
13/     2015/   240/    2946/  4529/   584/    723/   devices        sys/
1306/   2018/   241/    2947/  453/    5858/   7236/   diskstats      sysrq-trigger
14/     2041/   242/    2950/  472/    5872/   725/    dma            sysvipc/
1405/   2046/   243/    2951/  473/    5878/   726/    driver/        thread-self@
1449/   2049/   244/    2952/  476/    593/    728/   execdomains    timer_list
1457/   2055/   245/    2953/  477/    594/    7312/   fb             timer_stats
1470/   2059/   246/    2954/  479/    596/    7313/   filesystems    tty/
1490/   2062/   24697/  2955/  480/    597/    7321/   fs/            uptime
1495/   2070/   247/    2956/  481/    6130/   738/   interrupts     version
1508/   2082/   248/    2957/  482/    6131/   740/   iomem          vmallocinfo
1550/   2091/   249/    2965/  485/    616/    745/   ioports        vmnet/
1560/   2096/   24962/  2966/  486/    617/    746/   irq/           vmstat
1561/   2099/   2503/   3/     491/    6181/   748/   kallsyms       zoneinfo
1587/   21/     2506/   30/    497/    624/    749/   kcore
16/     2111/   2513/   3072/  498/    625/    752/   keys
1626/   2117/   2514/   3079/  499/    627/    758/   key-users
1664/   2120/   2516/   3090/  5/     628/    759/   kmsg
1669/   2125/   2517/   31/    501/    631/    762/   kpagecount
1675/   2137/   2520/   32/    502/    632/    763/   kpageflags
```



```
1685/ 2173/ 2521/ 3256/ 504/ 634/ 765/ latency_stats
1698/ 2175/ 2523/ 33/ 507/ 635/ 767/ loadavg
```

Notice many of the directory names are numbers; each corresponds to a running process and the name is the **process ID**. An important subdirectory we will discuss later is `/proc/sys`, under which many system parameters can be examined or modified.

2. View the following files:

- `/proc/cpuinfo`:
- `/proc/meminfo`:
- `/proc/mounts`:
- `/proc/swaps`:
- `/proc/version`:
- `/proc/partitions`:
- `/proc/interrupts`:

The names give a pretty good idea about what information they reveal.

Note that this information is not being constantly updated; it is obtained only when one wants to look at it.

3. Take a peek at any random process directory (if it is not a process you own some of the information might be limited unless you use **sudo**):

```
$ ls -F 5564
auxv          cwd@      latency   net/      projid_map statm
cgroup        environ  limits    ns/       root@     status
clear_refs    exe@     maps      oom_adj   sched     syscall
cmdline       fd/      mem       oom_score schedstat task/
comm          fdinfo/  mountinfo oom_score_adj smaps     uid_map
coredump_filter gid_map  mounts    pagemap   stack     wchan
cpuset        io       mountstats personality stat
```

Take a look at some of the fields in here such as: `cmdline`, `cwd`, `environ`, `mem`, and `status`

Chapter 6

Kernel Services and Configuration



Lab 6.1: System Tunables with sysctl

1. Check if you can **ping** your own system. (Note on **RHEL 7** you must be root to run **ping** on most external network addresses.)
2. Check the current value of `net.ipv4.icmp_echo_ignore_all`, which is used to turn on and off whether your system will respond to **ping**. A value of `0` allows your system to respond to pings.
3. Set the value to `1` using the **sysctl** command line utility and then check if pings are responded to.
4. Set the value back to `0` and show the original behavior is restored.
5. Now change the value by modifying `/etc/sysctl.conf` and force the system to activate this setting file without a reboot.
6. Check that this worked properly.

You will probably want to reset your system to have its original behavior when you are done.

Solution 6.1

You can use either `localhost`, `127.0.0.1` (loopback address) or your actual IP address for target of **ping** below.

1. `$ ping localhost`
2. `$ sysctl net.ipv4.icmp_echo_ignore_all`
3. `$ sudo sysctl net.ipv4.icmp_echo_ignore_all=1`
`$ ping localhost`

4.

```
$ sudo sysctl net.ipv4.icmp_echo_ignore_all=0
```



```
$ ping localhost
```
5. Add the following line to `/etc/sysctl.conf`:

```
net.ipv4.icmp_echo_ignore_all=1
```


and then do:

```
$ sysctl -p
```
6.

```
$ sysctl net.ipv4.icmp_echo_ignore_all
```



```
$ ping localhost
```

Since the changes to `/etc/sysctl.conf` are persistent, you probably want to restore things to its previous state.

Lab 6.2: Changing the Maximum Process ID

The normal behavior of a **Linux** system is that process IDs start out at `PID=1` for the `init` process, the first user process on the system, and then go up sequentially as new processes are constantly created (and die as well.)

However, when the PID reaches the value shown `/proc/sys/kernel/pid_max`, which is conventionally 32768 (32K), they will wrap around to lower numbers. If nothing else, this means you can't have more than 32K processes on the system since there are only that many slots for PIDs.

1. Obtain the current maximum PID value.
2. Find out what current PIDs are being issued
3. Reset `pid_max` to a lower value than the ones currently being issued.
4. Start a new process and see what it gets as a PID.

Solution 6.2

In the below we are going to use two methods, one involving `sysctl`, the other directly echoing values to `/proc/sys/kernel/pid_max`. Note that the `echo` method requires you to be root; `sudo` won't work. We'll leave it to you to figure out why, if you don't already know!

1.

```
$ sysctl kernel.pid_max
```



```
$ cat /proc/sys/kernel/pid_max
```
2. Type:

```
$ cat &
```



```
[1] 29222
```



```
$ kill -9 29222
```
3.

```
$ sudo sysctl kernel.pid_max=24000
```



```
$ echo 24000 > /proc/sys/kernel/pid_max # This must be done as root
```



```
$ cat /proc/sys/kernel/pid_max
```
4.

```
$ cat &
```



```
[2] 311
```



```
$ kill -9 311
```

Note that when starting over, the kernel begins at `PID=300`, not a lower value. You might notice that assigning PIDs to new processes is actually not trivial; since the system may have already turned over, the kernel always has to check when generating new PIDs that the PID is not already in use. The **Linux** kernel has a very efficient way of doing this that does not depend on the number of processes on the system.

Chapter 7

Kernel Modules



Lab 7.1: Kernel Modules

1. List all currently loaded kernel modules on your system.
2. Load a currently unloaded module on your system .

If you are running a distribution kernel, this is easy to find; you can simply look in the `/lib/modules/<kernel-version>/kernel/drivers/net` directory and grab one. (Distribution kernels come with drivers for every device, filesystem, network protocol etc. that a system might need.) However, if you are running a custom kernel you may not have many unloaded modules compiled.

3. Re-list all loaded kernel modules and see if your module was indeed loaded.
4. Remove the loaded module from your system.
5. Re-list again and see if your module was properly removed.

Solution 7.1

1. `$ lsmod`
2. In the following, substitute whatever module name you used for `3c59x`. Either of these methods work but, of course, the second is easier.
`$ sudo insmod /lib/modules/$(uname -r)/kernel/drivers/net/3c59.x`
`$ sudo /sbin/modprobe 3c59x`
3. `$ lsmod | grep 3c59x`
4. Once again, either method works.

```
$ sudo rmmod 3c59x  
$ sudo modprobe -r 3c59x
```

```
5. $ lsmod | grep 3c59x
```

Chapter 8

Devices and udev



Lab 8.1: udev

1. Create and implement a rule on your system that will create a symlink called `myusb` when a **USB** device is plugged in.
2. Plug in a **USB** device to your system. It can be a pendrive, mouse, webcam, etc.

Note: If you are running a virtual machine under a hypervisor, you will have to make sure the **USB** device is seen by the guest, which usually is just a mouse click which also disconnects it from the host.
3. Get a listing of the `/dev` directory and see if your symlink was created.
4. Remove the **USB** device. (If it is a drive you should always **umount** it first for safety.)
5. See if your symbolic link still exists in `/dev`.

Solution 8.1

1. Create a file named `/etc/udev/rules.d/75-myusb.rules` and have it include just one line of content:

```
$ cat /etc/udev/rules.d/75-myusb.rules  
  
SUBSYSTEM=="usb", SYMLINK+="myusb"
```

Do not use the deprecated key value `BUS` in place of `SUBSYSTEM`, as recent versions of **udev** have removed it.

Note the name of this file really does not matter. If there was an `ACTION` component to the rule the system would execute it; look at other rules for examples.

2. Plug in a device.
3. `$ ls -lF /dev | grep myusb`

4. If the device has been moun ted:

```
$ umount /media/whatever
```

where `/media/whatever` is the mount point. Safely remove the device.

5.

```
$ ls -lF /dev | grep myusb
```


Chapter 9

Partitioning and Formatting Disks



Lab 9.1: Using a File as a Disk Partition Image

For the purposes of the exercises in this course you will need unpartitioned disk space. It need not be large, certainly one or two GB will suffice.

If you are using your own machine, you either have it or you don't. If you don't, you will have to shrink a partition and the filesystem on it (first!) and then make it available, using **gparted** and/or the steps we have outlined or will outline.

Or you can use the **loop device** mechanism with or without the **parted** program, as we will do in the first two exercises in this section.

If you have real physical unpartitioned disk space you do not **need** to do the following procedures, but it is still a very useful learning exercise.

We are going to create a file that will be used as a container for a full hard disk partition image, and for all intents and purposes can be used like a real hard partition. In the next exercise we will show how to put more than one partition on it and have it behave as an entire disk.

1. Create a file full of zeros 1 GB in length:

```
$ dd if=/dev/zero of=imagefile bs=1M count=1024
```

You can make a much smaller file if you like or don't have that much available space in the partition you are creating the file on.

2. Put a filesystem on it:

```
$ mkfs.ext4 imagefile
mke2fs 1.42.9 (28-Dec-2013)
imagefile is not a block special device.
Proceed anyway? (y,n) y
Discarding device blocks: done
.....
```

Of course you can format with a different filesystem, doing `mkfs.ext3`, `mkfs.vfat`, `mkfs.xfs` etc.

3. Mount it somewhere:

```
$ mkdir mntpoint
$ sudo mount -o loop imagefile mntpoint
```

You can now use this to your heart's content, putting files etc. on it.

4. When you are done unmount it with:

```
$ sudo umount mntpoint
```

An alternative method to using the `loop` option to mount would be:

```
$ sudo losetup /dev/loop2 imagefile
$ sudo mount /dev/loop2 mntpoint
....
$ sudo umount mntpoint
$ sudo losetup -d /dev/loop2
```

We'll discuss `losetup` in a subsequent exercise, and you can use `/dev/loop[0-7]` but you have to be careful they are not already in use, as we will explain.

You should note that using a loop device file instead of a real partition can be useful, but it is pretty worthless for doing any kind of measurements or benchmarking. This is because you are placing one filesystem layer on top of another, which can only have a negative effect on performance, and mostly you just use the behavior of the underlying filesystem the image file is created on.

Lab 9.2: Partitioning a Disk Image File

The next level of complication is to divide the container file into multiple partitions, each of which can be used to hold a filesystem, or a swap area.

You can reuse the image file created in the previous exercise or create a new one.

1. Run `fdisk` on your imagefile:

```
$ sudo fdisk -C 130 imagefile
Device does not contain a recognized partition table
Building a new DOS disklabel with disk identifier 0x6280ced3.
Welcome to fdisk (util-linux 2.23.2).

Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.

Command (m for help):
```

2. Type `m` to get a list of commands:

```
Command (m for help): m

Command action
a   toggle a bootable flag
b   edit bsd disklabel
c   toggle the dos compatibility flag
d   delete a partition
g   create a new empty GPT partition table
G   create an IRIX (SGI) partition table
l   list known partition types
m   print this menu
```

```

n  add a new partition
o  create a new empty DOS partition table
p  print the partition table
q  quit without saving changes
s  create a new empty Sun disklabel
t  change a partition's system id
u  change display/entry units
v  verify the partition table
w  write table to disk and exit
x  extra functionality (experts only)

```

Command (m for help):

- The `-C 130` which sets the number of phony cylinders in the drive is only necessary in old versions of `fdisk`, which unfortunately you will find on **RHEL 6**. However, it will do no harm on other distributions.
Create a new primary partition and make it 256 MB (or whatever size you would like):

```

Command (m for help): n
Partition type:
  p   primary (0 primary, 0 extended, 4 free)
  e   extended
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-2097151, default 2048):
Using default value 2048
Last sector, +sectors or +size{K,M,G} (2048-2097151, default 2097151): +256M
Partition 1 of type Linux and of size 256 MiB is set

```

- Add a second primary partition also of 256 MB in size:

```

Command (m for help): n
Partition type:
  p   primary (1 primary, 0 extended, 3 free)
  e   extended
Select (default p): p
Partition number (2-4, default 2): 2
First sector (526336-2097151, default 526336):
Using default value 526336
Last sector, +sectors or +size{K,M,G} (526336-2097151, default 2097151): +256M
Partition 2 of type Linux and of size 256 MiB is set

```

Command (m for help): p

```

Disk imagefile: 1073 MB, 1073741824 bytes, 2097152 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk label type: dos
Disk identifier: 0x6280ced3

```

	Device	Boot	Start	End	Blocks	Id	System
imagefile1			2048	526335	262144	83	Linux
imagefile2			526336	1050623	262144	83	Linux

- Write the partition table to disk and exit:

```

Command (m for help): w
The partition table has been altered!

Syncing disks.

```

While this has given us some good practice, we haven't yet seen a way to use the two partitions we just created. We'll start over in the next exercise with a method that lets us do so.

Lab 9.3: Using losetup and parted

We are going to experiment more with:

- Loop devices and **losetup**
- **parted** to partition at the command line non-interactively.

We expect that you should read the **man pages** for **losetup** and **parted** before doing the following procedures.

Once again, you can reuse the image file or, better still, zero it out and start freshly or with another file.

1. Associate the image file with a **loop** device:

```
$ sudo losetup -f
/dev/loop1
$ sudo losetup /dev/loop1 imagefile
```

where the first command finds the first **free** loop device. The reason to do this is you may already be using one or more loop devices. For example, on the system that this is being written on, before the above command is executed:

```
$ losetup -a
/dev/loop0: []: (/usr/src/KERNELS.sqfs)
```

a **squashfs** compressed, read-only filesystem is already mounted using **/dev/loop0**. (The output of this command will vary with distribution.) If we were to ignore this and use **losetup** on **/dev/loop0** we would almost definitely corrupt the file.

2. Create a disk partition label on the loop device (image file):

```
$ sudo parted -s /dev/loop1 mklabel msdos
```

3. Create three primary partitions on the loop device:

```
$ sudo parted -s /dev/loop1 unit MB mkpart primary ext4 0 256
$ sudo parted -s /dev/loop1 unit MB mkpart primary ext4 256 512
$ sudo parted -s /dev/loop1 unit MB mkpart primary ext4 512 1024
```

4. Check the partition table:

```
$ fdisk -l /dev/loop1
Disk /dev/loop1: 1073 MB, 1073741824 bytes, 2097152 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk label type: dos
Disk identifier: 0x00050c11

   Device Boot      Start         End      Blocks   Id  System
/dev/loop1p1        1         500000       250000    83  Linux
/dev/loop1p2     500001     1000000       250000    83  Linux
/dev/loop1p3     1000001     2000000       500000    83  Linux
```

5. What happens next depends on what distribution you are on. For example, on **RHEL 7** and **Ubuntu 14.04** you will find new device nodes have been created:

```
$ ls -l /dev/loop1*
brw-rw---- 1 root disk  7, 1 Oct  7 14:54 /dev/loop1
brw-rw---- 1 root disk 259, 0 Oct  7 14:54 /dev/loop1p1
brw-rw---- 1 root disk 259, 3 Oct  7 14:54 /dev/loop1p2
brw-rw---- 1 root disk 259, 4 Oct  7 14:54 /dev/loop1p3
```

and we will use them in the following. However, on **RHEL 6** such nodes do not appear. Instead, you have to do:

```
$ sudo kpartx -lv /dev/loop1
$ sudo kpartx -av /dev/loop1
$ ls -l /dev/mapper/loop1*
lrwxrwxrwx 1 root root 7 Oct  9 07:12 /dev/mapper/loop1p1 -> ../dm-8
lrwxrwxrwx 1 root root 7 Oct  9 07:12 /dev/mapper/loop1p2 -> ../dm-9
lrwxrwxrwx 1 root root 8 Oct  9 07:12 /dev/mapper/loop1p3 -> ../dm-10
```

to associate device nodes with the partitions. So in what follows you can replace `/dev/loop1p[1-3]` with the actual names under `/dev/mapper`, or even easier you can do:

```
$ sudo ln -s /dev/mapper/loop1p1 /dev/loop1p1
$ sudo ln -s /dev/mapper/loop1p2 /dev/loop1p2
$ sudo ln -s /dev/mapper/loop1p3 /dev/loop1p3
```

6. Put filesystems on the partitions :

```
$ sudo mkfs.ext3 /dev/loop1p1
$ sudo mkfs.ext4 /dev/loop1p2
$ sudo mkfs.vfat /dev/loop1p3
```

7. Mount all three filesystems and show they are available:

```
$ mkdir mnt1 mnt2 mnt3

$ sudo mount /dev/loop1p1 mnt1
$ sudo mount /dev/loop1p2 mnt2
$ sudo mount /dev/loop1p3 mnt3

$ df -Th
Filesystem      Type      Size  Used Avail Use% Mounted on
/dev/sda1       ext4      29G   8.5G   19G  32% /
...
/dev/loop1p1    ext3      233M   2.1M  219M   1% mnt1
/dev/loop1p2    ext4      233M   2.1M  215M   1% mnt2
/dev/loop1p3    vfat      489M   0     489M   0% mnt3
```

8. After using the filesystems to your heart's content you can unwind it all:

```
$ sudo umount mnt1 mnt2 mnt3
$ rmdir mnt1 mnt2 mnt3
$ sudo losetup -d /dev/loop0
```

Lab 9.4: Partitioning a Real Hard Disk

If you have real hard disk un-partitioned space available, experiment with **fdisk** to create new partitions, either primary or logical within an extended partition. Write the new partition table to disk and then format and mount the new partitions.

Chapter 10

Encrypting Disks



Lab 10.1: Disk Encryption

In this exercise, you will encrypt a partition on the disk in order to provide a measure of security in the event that the hard drive or laptop is stolen. Reviewing the **cryptsetup** documentation first would be a good idea (`man cryptsetup` and `cryptsetup --help`).

1. Create a new partition for the encrypted block device with **fdisk**. Make sure the kernel is aware of the new partition table. A reboot will do this but there are other methods.
2. Format the partition with **cryptsetup** using **LUKS** for the crypto layer.
3. Create the un-encrypted pass through device by opening the crypted block device , i.e., **secret-disk**.
4. Add an entry to `/etc/crypttab` so that the system prompts for the passphrase on reboot.
5. Format the filesystem as an **ext4** filesystem.
6. Create a mount point for the new filesystem, ie. `/secret`.
7. Add an entry to `/etc/fstab` so that the filesystem is mounted on boot.
8. Try and mount the encrypted filesystem.
9. Validate the entire configuration by rebooting.

Solution 10.1

1. `$ sudo fdisk /dev/sda`

Create a new partition (in the below `/dev/sda4` to be concrete) and then either issue:

```
$ sudo partprobe -s
```

to have the system re-read the modified partition table, or reboot (which is far safer).

Note: If you can't use a real partition, use the technique in the previous chapter to use a loop device or image file for the same purpose.

2. `$ sudo cryptsetup luksFormat /dev/sda4`

3. `$ sudo cryptsetup luksOpen /dev/sda4 secret-disk`

4. Add the following to `/etc/crypttab`:

```
secret-disk    /dev/sda4
```

5. `$ sudo mkfs -t ext4 /dev/mapper/secret-disk`

6. `$ sudo mkdir -p /secret`

7. Add the following to `/etc/fstab`:

```
/dev/mapper/secret-disk    /secret    ext4    defaults    1 2
```

8. Mount just the one filesystem:

```
$ sudo mount /secret
```

or mount all filesystems mentioned in `/etc/fstab`:

```
$ sudo mount -a
```

9. Reboot.

Lab 10.2: Encrypted Swap

In this exercise, we will be encrypting the **swap partition**. Data written to the swap device can contain sensitive information. Because swap is backed by an actual partition, it is important to consider the security implications of having an unencrypted swap partition.

The process for encrypting is similar to the previous exercise, except we will not create a file system on the encrypted block device.

In this case, we are also going to use the existing swap device by first de-activating it and then formatting it for use as an encrypted swap device. It would be a little bit safer to use a fresh partition below, or you can safely reuse the encrypted partition you set up in the previous exercise. At the end we explain what to do if you have problems restoring.

(We will discuss swap management in a later chapter, but will show the few and easy commands for dealing with swap partitions here.)

You may want to revert back to the original unencrypted partition when we are done by just running **mkswap** on it again when it is not being used.

1. Find out what partition you are currently using for swap and then deactivate it:

```
$ cat /proc/swaps
Filename                                Type              Size  Used  Priority
/dev/sda11                             partition         41937760 -1
$ sudo swapoff /dev/sda11
```

2. Do the same steps as in the previous exercise to set up encryption:


```
$ sudo cryptsetup luksFormat /dev/sda11 # may use --cipher aes option
$ sudo cryptsetup luksOpen /dev/sda11 swapcrypt
```

3. Format the encrypted device to use with swap:

```
$ sudo mkswap /dev/mapper/swapcrypt
```

4. Now test to see if it actually works by activating it:

```
$ sudo swapon /dev/mapper/swapcrypt
$ cat /proc/swaps
```

5. To ensure the encrypted swap partition can be activated at boot you need to do two things:

- (a) Add a line to `/etc/crypttab` so that the system prompts for the passphrase on reboot:

```
swapcrypt /dev/sda11 /dev/urandom swap,cipher=aes-cbc-essiv:sha256,size=256
```

(Note `/dev/urandom` is preferred over `/dev/random` for reasons involving potential **entropy shortages** as discussed in the **man** page for **crypttab**.) You don't need the detailed options that follow, but we give them as an example of what more you can do.

- (b) Add an entry to the `/etc/fstab` file so that the swap device is activated on boot.

```
/dev/mapper/swapcrypt none swap defaults 0 0
```

6. You can validate the entire configuration by rebooting.

To restore your original unencrypted partition:

```
$ sudo swapoff /dev/mapper/swapcrypt
$ sudo cryptsetup luksClose swapcrypt
$ sudo mkswap /dev/sda11
$ sudo swapon -a
```

If the **swapon** command fails it is likely because `/etc/fstab` no longer properly describes the swap partition. If this partition is described in there by actual device node (`/dev/sda11`) there won't be a problem. You can fix either by changing the line in there to be:

```
/dev/sda11 swap swap defaults 0 0
```

or by giving a label when formatting and using it as in:

```
$ sudo mkswap -L SWAP /dev/sda11
```

and then putting in the file:

```
LABEL=SWAP swap swap defaults 0 0
```


Chapter 11

Linux Filesystems and the VFS



Lab 11.1: The tmpfs Special Filesystem

tmpfs is one of many special filesystems used under **Linux**. Some of these are not really used as filesystems, but just take advantage of the filesystem abstraction. However, **tmpfs** is a real filesystem that applications can do I/O on.

Essentially, **tmpfs** functions as a ramdisk; it resides purely in memory. But it has some nice properties that old-fashioned conventional ramdisk implementations did not have:

1. The filesystem adjusts its size (and thus the memory that is used) dynamically; it starts at zero and expands as necessary up to the maximum size it was mounted with.
2. If your RAM gets exhausted, **tmpfs** can utilize swap space. (You still can't try to put more in the filesystem than its maximum capacity allows, however.)
3. **tmpfs** does not require having a normal filesystem placed in it, such as **ext3** or **vfat**; it has its own methods for dealing with files and I/O that are aware that it is really just space in memory (it is not actually a block device), and as such are optimized for speed.

Thus there is no need to pre-format the filesystem with a **mkfs** command; you merely just have to mount it and use it.

Mount a new instance of **tmpfs** anywhere on your directory structure with a command like:

```
$ sudo mkdir /mnt/tmpfs
$ sudo mount -t tmpfs none /mnt/tmpfs
```

See how much space the filesystem has been given and how much it is using:

```
$ df -h /mnt/tmpfs
```

You should see it has been allotted a default value of half of your RAM; however, the usage is zero, and will only start to grow as you place files on `/mnt/tmpfs`.

You could change the allotted size as a mount option as in:

```
$ sudo mount -t tmpfs -o size=1G none /mnt/tmpfs
```

You might try filling it up until you reach full capacity and see what happens. Do not forget to unmount when you are done with:

```
$ sudo umount /mnt/tmpfs
```

Virtually all modern **Linux** distributions mount an instance of **tmpfs** at `/dev/shm`:

```
$ df -h /dev/shm
```

```
Filesystem      Type  Size  Used Avail Use% Mounted on
tmpfs           tmpfs 3.9G  24M  3.9G   1% /dev/shm
```

Many applications use this such as when they are using **POSIX** shared memory as an inter-process communication mechanism. Any user can create, read and write files in `/dev/shm`, so it is a good place to create temporary files in memory.

Create some files in `/dev/shm` and note how the filesystem is filling up with `df`.

In addition, many distributions mount multiple instances of **tmpfs**; for example, on a **RHEL 7** system:

```
$ df -h | grep tmpfs
```

```
devtmpfs          devtmpfs 3.9G    0 3.9G    0% /dev
tmpfs             tmpfs    3.9G  24M 3.9G    1% /dev/shm
tmpfs             tmpfs    3.9G  9.2M 3.9G    1% /run
tmpfs             tmpfs    3.9G    0 3.9G    0% /sys/fs/cgroup
/tmp/vmware-coop/564d9ea7-8e8e-29c0-2682-e5d3de3a51d8 tmpfs    3.3G    0 3.3G    0% /tmp/vmware-coop/
/tmp/vmware-coop/564d7668-ec55-ee45-f33e-c8e97e956190 tmpfs    2.3G  2.0G  256M   89% /tmp/vmware-coop/
none              tmpfs    1.0G  1.0G   0100% /tmp/ohno
```

Notice this was run on a system with 8 GB of ram, so clearly you can't have all these **tmpfs** filesystems actually using the 4 GB they have each been allotted!

Some distributions (such as **Fedora**) may (by default) mount `/tmp` as a **tmpfs** system; in such cases one has to avoid putting large files in `/tmp` to avoid running out of memory. Or one can disable this behavior as we discussed earlier when describing `/tmp`.

Chapter 12

Filesystem Features: Attributes, Creating, Checking, Mounting



Lab 12.1: Working with File Attributes

1. With your normal user account use **touch** to create an empty file named `/tmp/appendit`.
2. Use **cat** to append the contents of `/etc/hosts` to `/tmp/appendit`.
3. Compare the contents of `/tmp/appendit` with `/etc/hosts`; there should not be any differences.
4. Try to add the append-only attribute to `/tmp/appendit` by using **chattr**. You should see an error here. Why?
5. As root, retry adding the append-only attribute; this time it should work. Look at the file's extended attributes by using **lsattr**.
6. As a normal user, try and use **cat** to copy over the contents of `/etc/passwd` to `/tmp/appendit`. You should get an error. Why?
7. Try the same thing again as root. You should also get an error. Why?
8. As the normal user, again use the append redirection operator (`>>`) and try appending the `/etc/passwd` file to `/tmp/appendit`. This should work. Examine the resulting file to confirm.
9. As root, set the immutable attribute on `/tmp/appendit`, and look at the extended attributes again.
10. Try appending output to `/tmp/appendit`, try renaming the file, creating a hard link to the file, and deleting the file as both the normal user and as root.
11. We can remove this file by removing the extended attributes. Do so.

Solution 12.1

```

1. $ cd /tmp
   $ touch appendit
   $ ls -l appendit
   -rw-rw-r-- 1 coop coop 0 Oct 23 19:04 appendit

2. $ cat /etc/hosts > appendit

3. $ diff /etc/hosts appendit

4. $ chattr +a appendit
   chattr: Operation not permitted while setting flags on appendit

5. $ sudo chattr +a appendit
   $ lsattr appendit
   -----a-----e-- appendit

6. $ cat /etc/passwd > appendit
   bash: appendit: Operation not permitted

7. $ sudo su
   $ cat /etc/passwd > appendit
   bash: appendit: Operation not permitted
   $ exit

8. $ cat /etc/passwd >> /tmp/appendit
   $ cat appendit

9. $ sudo chattr +i appendit
   $ lsattr appendit
   ----ia-----e-- appendit

10. $ echo hello >> appendit
    -bash: appendit: Permission denied
    $ mv appendit appendit.rename
    mv: cannot move 'appendit' to 'appendit.rename': Operation not permitted
    $ ln appendit appendit.hardlink
    ln: creating hard link 'appendit.hardlink' => 'appendit': Operation not permitted
    $ rm -f appendit
    rm: cannot remove 'appendit': Operation not permitted

    $ sudo su
    $ echo hello >> appendit
    -bash: appendit: Permission denied
    $ mv appendit appendit.rename
    mv: cannot move 'appendit' to 'appendit.rename': Operation not permitted
    $ ln appendit appendit.hardlink
    ln: creating hard link 'appendit.hardlink' => 'appendit': Operation not permitted
    $ rm -f appendit
    rm: cannot remove 'appendit': Operation not permitted
    $ exit

11. $ sudo su
    $ lsattr appendit
    ----ia-----e-- appendit
    $ chattr -ia /appendit
    $ rm appendit
    rm: remove regular file 'appendit'? y
    $ ls appendit
    ls: cannot access appendit: No such file or directory

```

Lab 12.2: Mounting Options

In this exercise you will need to either create a fresh partition, or use a loopback file. The solution will differ slightly and we will provide details of both methods.

1. Use **fdisk** to create a new 250 MB partition on your system, probably on `/dev/sda`. Or create a file full of zeros to use as a loopback file to simulate a new partition.
2. Use **mkfs** to format a new filesystem on the partition or loopback file just created. Do this three times, changing the block size each time. Note the locations of the superblocks, the number of block groups and any other pertinent information, for each case.
3. Create a new subdirectory (say `/mnt/temppdir`) and mount the new filesystem at this location. Verify it has been mounted.
4. Unmount the new filesystem, and then remount it as read-only.
5. Try to create a file in the mounted directory. You should get an error here, why?
6. Unmount the filesystem again.
7. Add a line to your `/etc/fstab` file so that the filesystem will be mounted at boot time.
8. Mount the filesystem.
9. Modify the configuration for the new filesystem so that binary files may not be executed from the filesystem (change defaults to `noexec` in the `/mnt/temppdir` entry). Then remount the filesystem and copy an executable file (such as `/bin/ls`) to `/mnt/temppdir` and try to run it. You should get an error: why?

When you are done you will probably want to clean up by removing the entry from `/etc/fstab`.

Solution 12.2

Physical Partition Solution

1. We won't show the detailed steps in **fdisk**, as it is all ground covered earlier. We will assume the partition created is `/dev/sda11`, just to have something to show.

```
$ sudo fdisk /dev/sda
.....
w
$ partprobe -s
```

Sometimes the **partprobe** won't work, and to be sure the system knows about the new partition you have to reboot.

2.

```
$ sudo mkfs -t ext4 -v /dev/sda11
$ sudo mkfs -t ext4 -b 2048 -v /dev/sda11
$ sudo mkfs -t ext4 -b 4096 -v /dev/sda11
```

Note the `-v` flag (verbose) will give the requested information; you will see that for a small partition like this the default is 1024 byte blocks.

3.

```
$ sudo mkdir /mnt/temppdir
$ sudo mount /dev/sda11 /mnt/temppdir
$ mount | grep temppdir
```
4.

```
$ sudo umount /mnt/temppdir
$ sudo mount -o ro /dev/sda11 /mnt/temppdir
```

If you get an error while unmounting, make sure you are not currently in the directory.

5. `$ sudo touch /mnt/tempdir/afile`
6. `$ sudo umount /mnt/tempdir`
7. Put this line in `/etc/fstab`:
`/dev/sda11 /mnt/tempdir ext4 defaults 1 3`
8. `$ sudo mount /mnt/tempdir`
`$ sudo mount | grep tempdir`
9. Change the line in `/etc/fstab` to:
`/dev/sda11 /mnt/tempdir ext4 noexec 1 3`
 Then do:
`$ sudo mount -o remount /mnt/tempdir`
`$ sudo cp /bin/ls /mnt/tempdir`
`$ /mnt/tempdir/ls`

You should get an error here, why?

Loopback File Solution

1. `$ dd if=/dev/zero of=/tmp/imagefile bs=1M count=250`
2. `$ sudo mkfs -t ext4 -v /tmp/imagefile`
`$ sudo mkfs -t ext4 -b 2048 -v /tmp/imagefile`
`$ sudo mkfs -t ext4 -b 4096 -v /tmp/imagefile`

You will get warned that this is a file and not a partition, just proceed.

Note the `-v` flag (verbose) will give the requested information; you will see that for a small partition like this the default is 1024 byte blocks.

3. `$ sudo mkdir /mnt/tempdir`
`$ mount | grep tempdir`
4. `$ sudo umount /mnt/tempdir`
`$ sudo mount -o ro,loop /tmp/imagefile /mnt/tempdir`

If you get an error while unmounting, make sure you are not currently in the directory.

5. `$ sudo touch /mnt/tempdir/afile`
6. `$ sudo umount /mnt/tempdir`
7. Put this line in `/etc/fstab`:
`/tmp/imagefile /mnt/tempdir ext4 loop 1 3`
8. `$ sudo mount /mnt/tempdir`
`$ sudo mount | grep tempdir`
9. Change the line in `/etc/fstab` to:

`/tmp/imagefile /mnt/tempdir ext4 loop,noexec 1 3`
 Then do:
`$ sudo mount -o remount /mnt/tempdir`
`$ sudo cp /bin/ls /mnt/tempdir`
`$ /mnt/tempdir/ls`

You should get an error here, why?

Chapter 13

Filesystem Features: Swap, Quotas, Usage



Lab 13.1: Managing Swap Space

Examine your current swap space by doing:

```
$ cat /proc/swaps
```

Filename	Type	Size	Used	Priority
/dev/sda11	partition	41937760		-1

We will now add more swap space by adding either a new partition or a file. To use a file we can do:

```
$ dd if=/dev/zero of=swpfile bs=1M count=1024
```

```
1024+0 records in
1024+0 records out
1073741824 bytes (1.1 GB) copied, 1.30576 s, 822 MB/s
```

```
$ mkswap swpfile
```

```
Setting up swapspace version 1, size = 1048572 KiB
no label, UUID=85bb62e5-84b0-4fdd-848b-4f8a289f0c4c
```

(For a real partition just feed **mkswap** the partition name, but be aware all data on it will be erased!)
Activate the new swap space:

```
$ sudo swapon swpfile
```

```
swapon: /tmp/swpfile: insecure permissions 0664, 0600 suggested.
swapon: /tmp/swpfile: insecure file owner 500, 0 (root) suggested.
```

Notice **RHEL 7** warns us we are being insecure, we really should fix with:

```
$ sudo chown root:root swpfile
$ sudo chmod 600 swpfile
```

and ensure it is being used:

```
$ cat /proc/swaps

Filename                                Type              Size    Used    Priority
/dev/sda11                             partition         41937760      -1
/tmp/swpfile                           file             10485720      -2
```

Note the **Priority** field; swap partitions or files of lower priority will not be used until higher priority ones are filled.

Remove the swap file from use and delete it to save space:

```
$ sudo swapoff swpfile
$ sudo rm swpfile
```

Lab 13.2: Filesystem Quotas

1. Change the entry in `/etc/fstab` for your new filesystem to use user quotas (change `noexec` to `usrquota` in the entry for `/mnt/tempdir`). Then remount the filesystem.
2. Initialize quotas on the new filesystem, and then turn the quota checking system on.
3. Now set some quota limits for the normal user account: a soft limit of 500 blocks and a hard limit of 1000 blocks.
4. As the normal user, attempt to use `dd` to create some files to exceed the quota limits. Create `bigfile1` (200 blocks) and `bigfile2` (400 blocks).
You should get a warning. Why?
5. Create `bigfile3` (600 blocks).
You should get an error message. Why? Look closely at the file sizes.
6. Eliminate the persistent mount line you inserted in `/etc/fstab`.

Solution 13.2

1. Change `/etc/fstab` to have one of the following two lines according to whether you are using a real partition or a loopback file:

```
/dev/sda11    /mnt/tempdir ext4 usrquota    1 3
/tmp/imagefile /mnt/tempdir ext4 loop,usrquota 1 3
```

Then remount:

```
$ sudo mount -o remount /mnt/tempdir
```

2.

```
$ sudo quotacheck -u /mnt/tempdir
$ sudo quotaon -u /mnt/tempdir
$ sudo chown student.student /mnt/tempdir
```

(You won't normally do the line above, but we are doing it to make the next part easier).

3. Substitute your user name for the `student` user account.

4. `$ sudo edquota -u student`

```
5. $ cd /mnt/tempdir
$ dd if=/dev/zero of=bigfile1 bs=1024 count=200
200+0 records in
200+0 records out
204800 bytes (205 kB) copied, 0.000349604 s, 586 MB/s
$ quota
Disk quotas for user student (uid 500):
Filesystem blocks quota lim grace files qu lim gr
/dev/sda11  200  500 1000  1  0  0

$ dd if=/dev/zero of=bigfile2 bs=1024 count=400
sda11: warning, user block quota exceeded.
400+0 records in
400+0 records out
409600 bytes (410 kB) copied, 0.000654847 s, 625 MB/s
```

Create `bigfile3` (600 blocks).

```
6. $ quota
Disk quotas for user student (uid 500):
Filesystem blocks quota limit grace files qu lim gr
/dev/sda11  600*  500 1000 6days  2  0  0

$ dd if=/dev/zero of=bigfile3 bs=1024 count=600
sda11: write failed, user block limit reached.
dd: writing 'bigfile3': Disk quota exceeded
401+0 records in
400+0 records out
409600 bytes (410 kB) copied, 0.00177744 s, 230 MB/s
$ quota
Disk quotas for user student (uid 500):
Filesystem blocks quota limit grace files quota limit grace
/dev/sda11  1000*  500 1000 6days  3  0  0

$ ls -l
total 1068
-rw----- 1 root  root    7168 Dec 10 18:56 aquota.user
-rw-rw-r-- 1 student student 204800 Dec 10 18:58 bigfile1
-rw-rw-r-- 1 student student 409600 Dec 10 18:58 bigfile2
-rw-rw-r-- 1 student student 409600 Dec 10 19:01 bigfile3
drwx----- 2 root  root    16384 Dec 10 18:47 lost+found
-rwxr-xr-x 1 root  root    41216 Dec 10 18:52 more
```

Look closely at the file sizes.

7. Get rid of the line in `/etc/fstab`.

Chapter 14

The Ext2/Ext3/Ext4 Filesystems



Lab 14.1: Defragmentation

Newcomers to **Linux** are often surprised at the lack of mention of filesystem **defragmentation** tools, since such programs are routinely used in the **Windows** world.

However, native filesystems in **UNIX**-type operating systems, including **Linux**, tend not to suffer serious problems with filesystem fragmentation.

This is primarily because they do not try to cram files onto the innermost disk regions where access times are faster. Instead, they spread free space out throughout the disk, so that when a file has to be created there is a much better chance that a region of free blocks big enough can be found to contain the entire file in either just one or a small number of pieces.

For modern hardware, the concept of innermost disk regions is obscured by the hardware anyway; and for **SSDs** defragmentation would actually shorten the lifespan of the storage media due to finite read/erase/write cycles.

Furthermore, the newer **journaling** filesystems (including **ext4**) work with **extents** (large contiguous regions) by design.

However, there does exist a tool for defragmenting **ext4** filesystems:

```
$ sudo e4defrag
```

```
Usage : e4defrag [-v] file...| directory...| device...
       : e4defrag -c file...| directory...| device...
```

e4defrag is part of the **e2fsprogs** package and should be on all modern **Linux** distributions, although it doesn't come with **RHEL 6** which is somewhat long in tooth.

The only two options are:

- **-v**: Be verbose.
- **-c**: Don't actually do anything, just analyze and report.

The argument can be:

- A file
- A directory
- An entire device

Examples:

```
$ sudo e4defrag -c /var/log
```

```
<Fragmented files>          now/best      s ize/ext
1. /var/log/lastlog          5/1         9KB
2. /var/log/sa/sa24          3/1        80KB
3. /var/log/rhsm/rhsm.log     2/1       142 KB
4. /var/log/messages         2/1     4590 KB
5. /var/log/Xorg.1.log.old    1/1        36 KB

Total/best extents          120/112
Average size per extent     220 KB
Fragmentationscore          1
[0-30 no problem: 31-55 a little bit fragmented: 56- needs defrag]
This directory (/var/log) does not need defragmentation.
Done.
```

```
$ sudo e4defrag /var/log
```

```
ext4 defragmentation for directory(/var/log)
[2/152]/var/log/Xorg.2.log:   100%  [ OK ]
[3/152]/var/log/Xorg.0.log.old: 100%  [ OK ]
[4/152]/var/log/messages-20141019.gz: 100%  [ OK ]
[5/152]/var/log/boot.log:     100%  [ OK ]

[7/152]/var/log/cups/page_log-20140924.gz: 100%  [ OK ]
[8/152]/var/log/cups/access_log-20141019.gz: 100%  [ OK ]
[9/152]/var/log/cups/access_log: 100%  [ OK ]
[10/152]/var/log/cups/error_log-20141018.gz: 100%  [ OK ]
[11/152]/var/log/cups/error_log-20141019.gz: 100%  [ OK ]
[12/152]/var/log/cups/access_log-20141018.gz: 100%  [ OK ]
[14/152]/var/log/cups/page_log-20141018.gz: 100%  [ OK ]
...
[152/152]/var/log/Xorg.1.log.old: 100%  [ OK ]

Success: [ 112/152 ]
Failure: [ 40/152 ]
```

Try running **e4defrag** on various files, directories, and entire devices, always trying with **-c** first.

You will generally find that **Linux** filesystems only tend to need defragmentation when they get very full, over 90 percent or so, or when they are small and have relatively large files, like when a **boot** partition is used.

Lab 14.2: Modifying Filesystem Parameters with tune2fs

We are going to fiddle with some properties of a formatted **ext4** filesystem. This does not require unmounting the filesystem first.

In the below you can work with an image file you create as in:

```
$ dd if=/dev/zero of=imagefile bs=1M count=1024
```

or you can substitute `/dev/sdaX` (using whatever partition the filesystem you want to modify is mounted on) for `imagefile`.

1. Using `dumpe2fs`, obtain information about the filesystem whose properties you want to adjust.
2. Ascertain the maximum mount count setting (after which a filesystem check will be forced) and modify it to have the value 30.
3. Set the `Check interval` (the amount of time after which a filesystem check is forced), to three weeks.
4. Calculate the percent age of blocks reserved, and then reset it to 10%.

Solution 14.2

```
1. $ dumpe2fs imagefile > dump_results

2. $ grep -i "Mount count" dump_results
Mount count:          0
Maximum mount count:  -1

$ sudo tune2fs -c 30 imagefile
$ grep -i "Mount count" dump_results
Mount count:          0
Maximum mount count:  30

3. $ grep -i "Check interval" dump_results
Check interval:        0 (<none>)

$ sudo tune2fs -i 3w imagefile
$ grep -i "Check interval" dump_results
Check interval:        1814400 (3 weeks)

4. $ grep -i "Block Count" dump_results
Block count:           131072
Reserved block count:  6553

$ echo "scale=4; 6553/131072" | bc
.0499

$ sudo tune2fs -m 10 imagefile
$ tune2fs 1.42.9 (28-Dec-2013)
Setting reserved blocks percentage to 10% (13107 blocks)
$ grep -i "Block Count" dump_results
Block count:           131072
Reserved block count:  13107
```


Chapter 15

The XFS and btrfs Filesystems



Lab 15.1: Finding Out More About xfs

We do not have a detailed lab exercise you can do with `xfs`; many systems still will not have the kernel modules and relevant user utilities installed. However, if your **Linux** kernel and distribution does support it, you can easily create a filesystem with `mkfs -t xfs`.

Then you can find out about available `xfs`-related utilities with:

```
$ man -k xfs
```

```
attr (1)          - extended attributes on XFS filesystem objects
filesystems (5)    - Linux file-system types: minix, ext, ext2, ext3, ext4,...
fs (5)            - Linux file-system types: minix, ext, ext2, ext3, ext4,...
fsck.xfs (8)       - do nothing, successfully
fsfreeze (8)       - suspend access to a filesystem (Linux Ext3/4, ReiserFS...)
mkfs.xfs (8)       - construct an XFS filesystem
pmdaxfs (1)        - XFS filesystem performance metrics domain agent (PMDA)
xfs (5)           - layout of the XFS filesystem
xfs_admin (8)      - change parameters of an XFS filesystem
xfs_bmap (8)       - print block mapping for an XFS file
xfs_copy (8)       - copy the contents of an XFS filesystem
xfs_db (8)         - debug an XFS filesystem
xfs_estimate (8)   - estimate the space that an XFS filesystem will take
xfs_freeze (8)     - suspend access to an XFS filesystem
xfs_fsr (8)        - filesystem reorganizer for XFS
xfs_growfs (8)     - expand an XFS filesystem
xfs_info (8)       - expand an XFS filesystem
xfs_io (8)         - debug the I/O path of an XFS filesystem
xfs_logprint (8)   - print the log of an XFS filesystem
xfs_mdrestore (8)  - restores an XFS met adump image to a filesystem image
xfs_metadump (8)   - copy XFS filesystem metadata to a file
xfs_mkfile (8)     - create an XFS file
xfs_ncheck (8)     - generate pathnames from i-numbers for XFS
```

```

xfs_quota (8)      - manage use of quota on XFS filesystems
xfs_repair (8)     - repair an XFS filesystem
xfs_rtcp (8)       - XFS realtime copy command
xfsdump (8)        - XFS filesystem incremental dump utility
xfsinvtutil (8)    - xfsdump inventory database checking and pruning utility
xfsrestore (8)     - XFS filesystem incremental restore utility
xqmstats (8)       - Display XFS quota manager statistics from /proc

```

Read about these utility programs and see if you can play with them on the filesystem you created.

Lab 15.2: Finding Out More About btrfs

We defer to a detailed lab exercise, however, if your Linux kernel and distribution support it, you can easily create a filesystem with `mkfs -t btrfs`.

Then you can find out about available **btrfs**-related utilities with:

```

$ man -k btrfs

btrfs-image (8)      - create/restore an image of the filesystem
btrfs-show (8)       - scan the /dev directory for btrfs partitions and print...
btrfsck (8)          - check a btrfs filesystem
btrfsctl (8)         - control a btrfs filesystem
mkfs.btrfs (8)       - create an btrfs filesystem
btrfs (8)            - control a btrfs filesystem
btrfs-convert (8)    - convert ext2/3/4 to btrfs.
btrfs-debug-tree (8) - dump Btrfs filesystem metadata into stdout.
btrfs-find-root (8)  - filter to find btrfs root.
btrfs-map-logical (8) - map btrfs logical extent to physical extent
btrfs-show-super (8) - show btrfs superblock information stored in devices
btrfs-zero-log (8)   - clear out log tree.

btrfstune (8)        - tune various filesystem parameters.

```

Read about these utility programs and see if you can play with them on the filesystem you created.

Chapter 16

Logical Volume Management (LVM)



Lab 16.1: Logical Volumes

We are going to create a logical volume using two 250 MB partitions. We are going to assume you have real partitionable disk space available.

1. Create two 250 MB partitions of type logical volume e (**8e**).
2. Convert the partitions to physical volumes.
3. Create a volume group named **myvg** and add the two physical volumes to it. Use the default extent size.
4. Allocate a 300 MB logical volume named **mylvm** from volume group **myvg**.
5. Format and mount the logical volume **mylvm** at **/mylvm**.
6. Use **lvdisplay** to view information about the logical volume.
7. Grow the logical volume and corresponding filesystem to 350 MB.

Solution 16.1

1. Execute:

```
$ sudo fdisk /dev/sda
```

using whatever hard disk is appropriate, and create the two partitions. While in **fdisk**, typing **t** will let you set the partition type to **8e**. While it doesn't matter if you don't set the type, it is a good idea to lessen confusion. Use **w** to rewrite the partition table and exit, and then

```
$ sudo partprobe -s
```

or reboot to make sure the new partitions take effect.

2. Assuming the new partitions are `/dev/sdaX` and `/dev/sdaY`:

```
$ sudo pvcreate /dev/sdaX
$ sudo pvcreate /dev/sdaY
$ sudo pvdisplay
```

3. `$ sudo vgcreate myvg /dev/sdaX /dev/sdaY`
`$ sudo vgdisplay`

4. `$ sudo lvcreate -L 300M -n mylvm myvg`
`$ sudo lvdisplay`

```
$ sudo mkfs.ext4 /dev/myvg/mylvm
```

5. `$ mkdir /mylvm`
`$ sudo mount /dev/myvg/mylvm /mylvm`

If you want the mount to be persistent, edit `/etc/fstab` to include the line:

```
/dev/myvg/mylvm /mylvm ext4 defaults 0 0
```

6. `$ sudo lvdisplay`

7. `$ df -h`
`$ sudo lvextend -L 350M /dev/myvg/mylvm`
`$ sudo resize2fs /dev/myvg/mylvm`
`$ df -h`

or

```
$ sudo lvextend -r -L +50M /dev/myvg/mylvm
```

Chapter 17

RAID



Lab 17.1: Creating a RAID Device

Normally when creating a **RAID** device we would use partitions on separate disks. However, for this exercise we probably don't have such hardware available.

(Note we will need to have two partitions on the same disk, or we can use **LVM** partitions just for demonstration purposes. (Note we can't use image files and loopback for this exercise.)

The process will be the same whether the partitions are on one drive or several (Although there is obviously little reason to actually create a **RAID** on a single device).

1. Create two 200 MB partitions of type raid (**fd**) either on your hard disk using **fdisk**, or using **LVM**.
2. Create a **RAID 1** device named **/dev/md0** using the two partitions.
3. Format the **RAID** device as an **ext4** filesystem. Then mount it at **/myraid** and make the mount persistent.
4. Place the information about **/dev/md0** in **/etc/mdadm.conf** file using **mdadm**. (Depending on your distribution, this file may not previously exist.)
5. Examine **/proc/mdstat** to see the status of your **RAID** device.

Solution 17.1

1. If you are using real hard disk part itions do

```
$ sudo fdisk /dev/sda
```

and create the partitions as we have done before. For purposes of being definite, we will call them **/dev/sdaX** and **/dev/sdaY**. You will need to run **partprobe** or **kpartx** or reboot after you are done to make sure the system is properly aware of the new partitions.

LVM partitions will be perfectly fine for this exercise and can be easily created with:

```
$ sudo lvcreate -L 200M -n MD1 VG
$ sudo lvcreate -L 200M -n MD2 VG
```

where we have assumed `VG` to be the name of the volume group. Nothing needs to be done after creation to make sure the system is aware of the new LVM partitions.

2. `$ sudo mdadm -C /dev/md0 --level=1 --raid-disks=2 /dev/sdaX /dev/sdaY`

or

```
$ sudo mdadm -C /dev/md0 --level=1 --raid-disks=2 /dev/VG/MD1 /dev/VG/MD2
```

3. `$ sudo mkfs.ext4 /dev/md0`
`$ sudo mkdir /myraid`
`$ sudo mount /dev/md0 /myraid`

and add to `/etc/fstab`

```
/dev/md0 /myraid ext4 defaults 0 0
```

4. `$ mdadm --detail --scan >> /etc/mdadm.conf`

5. `$ cat /proc/mdstat`
- ```
Personalities : [raid1]
md0 : active raid1 dm-14[1] dm-13[0]
 204736 blocks [2/2] [UU]

unused devices: <none>
```

You should probably verify that with a reboot, the **RAID** volume is mounted automatically. When you are done, you probably will want to clean up by removing the line from `/etc/fstab`, and then getting rid of the partitions.

## Chapter 18

# Local System Security



### Lab 18.1: Security and Mount Options

We are going to mount a partition or loop device with the `noexec` option to prevent execution of programs that reside on the filesystem therein. You can certainly do this with a pre-existing and mounted partition, but you may not be able to easily change the behavior while the partition is mounted. Therefore, to demonstrate we'll use a loop device, which is a harmless procedure.

1. Set up an empty file, put a filesystem on it and mount it.
2. Copy an executable file to it from somewhere else on your system and test that it works in the new location.
3. Unmount it and remount with the `noexec` option.
4. Test if the executable still works. It should give you an error because of the `noexec` mount option.
5. Clean up.

### Solution 18.1

1. 

```
$ dd if=/dev/zero of=image bs=1M count=100
$ sudo mkfs.ext3 image
$ mkdir mountpoint
$ sudo mount -o loop image mountpoint
```
2. 

```
$ sudo cp /bin/ls mountpoint
$ mountpoint/ls
```
3. 

```
$ sudo umount mountpoint
$ sudo mount -o noexec,loop image mountpoint
```

```

OR

$ sudo mount -o noexec,remount image mountpoint

4. $ mountpoint/ls

5. $ sudo umount mountpoint
 $ rm image
 $ rmdir mountpoint

```

Note that this is not persistent. To make it persistent you would need to add the option to `/etc/fstab` with a line like:

```
/home/student/image /home/student/mountpoint ext3 loop,rw,noexec 0 0
```

## Lab 18.2: More on setuid and Scripts

Suppose we have the following C program (`./writeit.c`) which attempts to overwrite a file in the current directory named `afile`:

```

#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <string.h>
#include <stdlib.h>
#include <sys/stat.h>

int main(int argc, char *argv[])
{
 int fd, rc;
 char *buffer = "TESTING A WRITE";
 fd = open("./afile", O_RDWR | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
 rc = write(fd, buffer, strlen(buffer));
 printf("wrote %d bytes\n", rc);
 close(fd);
 exit(EXIT_SUCCESS);
}

```

If you are taking the online self-paced version of this course, the source code is available for download from your **Lab** screen.

If the program is called `writeit.c`, it can be compiled simply by doing:

```
$ make writeit
```

or equivalently

```
$ gcc -o writeit writeit.c
```

If (as a normal user) you try to run this program on a file owned by root you'll get

```
$ sudo touch afile
$./writeit
```

```
wrote -1 bytes
```



but if you run it as root:

```
$ sudo ./writeit
wrote 15 bytes
```

Thus, the root user was able to overwrite the file it owned, but a normal user could not.

Note that changing the owner of `writeit` to root does not help:

```
$ sudo chown root.root writeit
$./writeit
wrote -1 bytes
```

because it still will not let you clobber `afile`.

By setting the `setuid` bit you can make any normal user capable of doing it:

```
$ sudo chmod +s writeit
$./writeit
wrote 15 bytes
```

You may be asking, why didn't we just write a script to do such an operation, rather than to write and compile an executable program?

Under **Linux**, if you change the `setuid` on such an executable script, it won't do anything unless you actually change the `setuid` bit on the shell (such as `bash`) which would be a big mistake; anything running from then on would have escalated privilege!



## Chapter 19

# Linux Security Modules



### Lab 19.1: SELinux

Before starting this exercise verify **SELinux** is enabled and in **enforcing** mode, by editing `/etc/selinux/config` and rebooting if necessary.

Obviously you can only do this on a system such as **RHEL** where **SELinux** is installed.

1. Install the **vsftpd** and **ftp** packages.
2. Create a user account **user1** with the password **password**.
3. Change to **user1** account and write some text to a file named `/home/user1/user1file`.
4. Exit the **user1** account and make sure the **ftp** (**vsftpd** by name) service is running.
5. **ftp** to **localhost**, login as **user1**, and try to get **user1file**. It should fail.  
Note this step can fail either at the login, or at the file transfer. The fix for both problems is the same, so it should not affect the exercise. This difference in the behavior is a consequence of differences in the **SELinux** policy.
6. Check `/var/log/messages` to see why. You should see an error from **setroubleshoot**. Run the **sealert** command shown earlier.
7. Fix the error, and now try to **ftp**, login as **user1**, and get **user1file** again. This time it should work.

### Solution 19.1

1. `$ sudo yum install vsftpd ftp`

2. `$ sudo useradd user1`  
`$ sudo passwd user1`  
Changing password for user user1.  
New password: password  
BAD PASSWORD: The password fails the dictionary check - it is based on a dictionary word  
Retype new password: password  
passwd: all authentication tokens updated successfully.
3. `$ sudo su - user1`  
[user1@rhel7 ~]\$ echo 'file created at /home/user1' > userfile  
[user1@rhel7 ~]\$ ls  
userfile
4. [user1@rhel7 ~]\$ exit  
`$ sudo systemctl status vsftpd.service`  
vsftpd.service - Vsftpd ftp daemon  
Loaded: loaded (/usr/lib/systemd/system/vsftpd.service; disabled)  
Active: active (running) since Fri 2014-11-21 14:08:14 CET; 32min ago  
...  
5. `$ ftp localhost`  
Trying ::1...  
Connected to localhost (::1).  
220 (vsFTPd 3.0.2)  
Name (localhost:peter): user1  
331 Please specify the password.  
Password: password  
230 Login successful.  
Remote system type is UNIX.  
Using binary mode to transfer files.  
ftp> get userfile  
local: userfile remote: userfile  
229 Entering Extended Passive Mode (|||35032|).  
550 Failed to open file.  
ftp> quit  
221 Goodbye.
6. `$ tail /var/log/messages`  
Nov 21 14:23:26 rhel7 setroubleshoot: SELinux is preventing /usr/sbin/vsftpd from read access on the file .  
For complete SELinux messages. run `sealert -l 7f8e5e6f-bcee-4c59-9cd1-72b90fb1f462`  
\*\*\*\*\* Plugin catchall\_boolean (47.5 confidence) suggests \*\*\*\*\*  
  
If you want to allow ftp to home dir  
Then you must tell SELinux about this by enabling the 'ftp\_home\_dir' boolean.  
  
Do  
`setsebool -P ftp_home_dir 1`  
  
Notice that the suggestion to fix the issue can be found at the log file, and it is not even necessary to run `sealert`.
7. `$ sudo setsebool -P ftp_home_dir 1`  
`$ ftp localhost`  
Trying ::1...  
Connected to localhost (::1).  
220 (vsFTPd 3.0.2)  
Name (localhost:peter): user1  
331 Please specify the password.  
Password:

```
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp> get userifile
local: userifile remote: userifile
229 Entering Extended Passive Mode (||||18769|).
150 Opening BINARY mode data connection for userifile (28 bytes).
226 Transfer complete.
28 bytes received in 4.2e-05 secs (666.67 Kbytes/sec)
ftp> quit
221 Goodbye.

$ cat userifile
file created at /home/user1
```



## Chapter 20

# Processes



### Lab 20.1: Controlling Processes with ulimit

Please do:

```
$ help ulimit
```

and read `/etc/security/limits.conf` before doing the following steps.

1. Start a new shell by typing **bash** (or opening a new terminal) so that your changes are only effective in the new shell. View the current limit on the number of open files and explicitly view the hard and soft limits.
2. Set the limit to the hard limit value and verify if it worked.
3. Set the hard limit to 2048 and verify it worked.
4. Try to set the limit back to the previous value. Did it work?

### Solution 20.1

```
1. $ bash
$ ulimit -n
1024
$ ulimit -S -n
1024
$ ulimit -H -n
4096
```

```

2. $ ulimit -n hard
 $ ulimit -n
 4096

3. $ ulimit -n 2048
 $ ulimit -n
 2048

4. $ ulimit -n 4096
 bash: ulimit: open files: cannot modify limit: Operation not permitted
 $ ulimit -n
 2048

```

You can't do this anymore!

Note that if we had chosen a different limit, such as stack size ( `-s` ) we could raise back up again as the hard limit is unlimited.

## Lab 20.2: Examining System V IPC Activity

**System V IPC** is a rather old method of **I**nter **P**rocess **C**ommunication that dates back to the early days of **UNIX**. It involves three mechanisms:

1. Shared Memory Segments
2. Semaphores
3. Message Queues

More modern programs tend to use **POSIX IPC** methods for all three of these mechanisms, but there are still plenty of **System V IPC** applications found in the wild.

To get an overall summary of **System V IPC** activity on your system, do:

```
$ ipcs
```

```

----- Message Queues -----
key msqid owner perms used-bytes messages

----- Shared Memory Segments -----
key shmid owner perms bytes nattch status
0x011147030 root 600 1000 6
0x00000000 98305 coop 600 4194304 2 dest
0x00000000 196610 coop 600 4194304 2 dest
0x00000000 23068675 coop 700 1138176 2 dest
0x00000000 23101444 coop 600 393216 2 dest
0x00000000 23134213 coop 600 524288 2 dest
0x00000000 24051718 coop 600 393216 2 dest
0x00000000 23756807 coop 600 524288 2 dest
0x00000000 24018952 coop 600 67108864 2 dest
0x00000000 23363593 coop 700 95408 2 dest
0x00000000 1441811 coop 600 2097152 2 dest

----- Semaphore Arrays -----
key semid owner perms nsems
0x00000000 98304 apache 600 1
0x00000000 131073 apache 600 1
0x00000000 163842 apache 600 1
0x00000000 196611 apache 600 1
0x00000000 229380 apache 600 1

```



Note almost all of the currently running shared memory segments have a key of 0 (also known as `IPC_PRIVATE`) which means they are only shared between processes in a parent/child relationship. Furthermore, all but one are marked for destruction when there are no further attachments.

One can gain further information about the processes that have created the segments and last attached to them with:

```
$ ipcs -p
```

```
----- Message Queues PIDs -----
msqid owner lspid lrpid

----- Shared Memory Creator/Last-op PIDs -----
shmids owner cpid lpid
0 root 1023 1023
98305 coop 2265 18780
196610 coop 2138 18775
23068675 coop 989 1663
23101444 coop 989 1663
23134213 coop 989 1663
24051718 coop 20573 1663
23756807 coop 10735 1663
24018952 coop 17875 1663
23363593 coop 989 1663
1441811 coop 2048 20573
```

Thus, by doing:

```
$ ps aux |grep -e 20573 -e 2048
```

```
coop 2048 5.3 3.7 1922996 305660 ? Rl Oct27 77:07 /usr/bin/gnome-shell
coop 20573 1.9 1.7 807944 141688 ? Sl 09:56 0:01 /usr/lib64/thunderbird/thunderbird
coop 20710 0.0 0.0 112652 2312 pts/0 S+ 09:57 0:00 grep --color=auto -e 20573 -e 2048
```

we see **thunderbird** is using a shared memory segment created by **gnome-shell**.

Perform these steps on your system and identify the various resources being used and by who. Are there any potential **leaks** (shared resources no longer being used by any active processes) on the system? For example, doing:

```
$ ipcs
```

```
----- Shared Memory Segments -----
key shmids owner perms bytes nattch status

0x00000000 622601 coop 600 2097152 2 dest
0x0000001a 13303818 coop 666 8196 0
```

shows a shared memory segment with no attachments and not marked for destruction. Thus it might persist forever, leaking memory if no subsequent process attaches to it.



## Chapter 21

# Signals



### Lab 21.1: Examining Signal Priorities and Execution

We give you a C program that includes a signal handler that can handle any signal. The handler avoids making any system calls (such as those that might occur while doing I/O).

```
/*
 * Examining Signal Priorities and Execution.
 *
 * The code herein is: Copyright the Linux Foundation, 2014
 * Author: J. Cooperstein
 *
 * This Copyright is retained for the purpose of protecting free
 * redistribution of source.
 *
 * This code is distributed under Version 2 of the GNU General Public
 * License, which you should have received with the source.
 */

#include <stdio.h>
#include <unistd.h>
#include <signal.h>
#include <stdlib.h>
#include <string.h>
#include <pthread.h>

#define NUMSIGS 64

/* prototypes of locally-defined signal handlers */

void (sig_handler) (int);

int sig_count[NUMSIGS + 1]; /* counter for signals received */
```

```

volatile static int line = 0;
volatile int signumbuf[6400], sigcountbuf[6400];

int main(int argc, char *argv[])
{
 sigset_t sigmask_new, sigmask_old;
 struct sigaction sigact, oldact;
 int signum, rc, i;
 pid_t pid;

 pid = getpid();

 /* block all possible signals */
 rc = sigfillset(&sigmask_new);
 rc = sigprocmask(SIG_SETMASK, &sigmask_new, &sigmask_old);

 /* Assign values to members of sigaction structures */
 memset(&sigact, 0, sizeof(struct sigaction));
 sigact.sa_handler = sig_handler; /* we use a pointer to a handler */
 sigact.sa_flags = 0; /* no flags */
 /* VERY IMPORTANT */
 sigact.sa_mask = sigmask_new; /* block signals in the handler itself */

 /*
 * Now, use sigaction to create references to local signal
 * handlers * and raise the signal to myself
 */

 printf
 ("Installing signal handler and Raising signal for signal number:\n\n");
 for (signum = 1; signum <= NUMSIGS; signum++) {
 if (signum == SIGKILL || signum == SIGSTOP || signum == 32
 || signum == 33) {
 printf(" --");
 continue;
 }
 sigaction(signum, &sigact, &oldact);
 /* send the signal 3 times! */
 rc = raise(signum);
 rc = raise(signum);
 rc = raise(signum);
 if (rc) {
 printf("Failed on Signal %d\n", signum);
 } else {
 printf("%4d", signum);
 if (signum % 16 == 0)
 printf("\n");
 }
 }
 fflush(stdout);

 /* restore original mask */
 rc = sigprocmask(SIG_SETMASK, &sigmask_old, NULL);

 printf("\nSignal Number(Times Processed)\n");
 printf("-----\n");
 for (i = 1; i <= NUMSIGS; i++) {
 printf("%4d:%3d ", i, sig_count[i]);
 if (i % 8 == 0)
 printf("\n");
 }
 printf("\n");
}

```

```

 printf("\nHistory: Signal Number(Count Processed)\n");
 printf("-----\n");
 for (i = 0; i < line; i++) {
 if (i % 8 == 0)
 printf("\n");
 printf("%4d(%1d)", signumbuf[i], sigcountbuf[i])
 }
 printf("\n");
 exit(EXIT_SUCCESS);
}

void sig_handler(int sig)
{
 sig_count[sig]++;
 signumbuf[line] = sig;
 sigcountbuf[line] = sig_count[sig];
 line++;
}

```

If you are taking the online self-paced version of this course, the source code is available for download from your **Lab** screen.

You will need to compile it and run it as in:

```

$ gcc -o signals signals.c
$./signals

```

When run, the program:

- Does not send the signals **SIGKILL** or **SIGSTOP**, which can not be handled and will always terminate a program.
- Stores the sequence of signals as they come in, and updates a counter array for each signal that indicates how many times the signal has been handled.
- Begins by suspending processing of all signals and then installs a new set of signal handlers for all signals.
- Sends every possible signal to itself multiple times and then unblocks signal handling and the queued up signal handlers will be called.
- Prints out statistics including:
  - The total number of times each signal was received.
  - The order in which the signals were received, noting each time the total number of times that signal had been received up to that point.

Note the following:

- If more than one of a given signal is **raised** while the process has blocked it, does the process **receive** it multiple times? Does the behavior of **real time** signals differ from normal signals?
- Are all signals received by the process, or are some handled before they reach it?
- What order are the signals received in?

One signal, **SIGCONT** (18 on **x86**) may not get through; can you figure out why?

#### Note:

On some **Linux** distributions signals 32 and 33 can not be blocked and will cause the program to fail. Even though system header files indicate **SIGRTMIN**=32, the command **kill -1** indicates **SIGRTMIN**=34.

Note that **POSIX** says one should use signal names, not numbers, which are allowed to be completely implementation dependent.

You should generally avoid sending these signals.



## Chapter 22

# System Monitoring



### Lab 22.1: Using stress

**stress** is a C language program written by Amos Waterland at the University of Oklahoma, licensed under the **GPL v2**. It is designed to place a configurable amount of stress by generating various kinds of workloads on the system.

If you are lucky, you can install **stress** directly from your distribution's packaging system. Otherwise, you can obtain the source from <http://people.seas.harvard.edu/~spw/stress>, and then compile and install by doing:

```
$ tar zxvf stress-1.0.4.tar.gz
$ cd stress-1.0.4
$./configure
$ make
$ sudo make install
```

There may exist pre-packaged downloadable binaries in the **.deb** and **.rpm** formats; see the home page for details and locations.

Once installed, you can do:

```
$ stress --help
```

for a quick list of options, or

```
$ info stress
```

for more detailed documentation.

As an example, the command:

```
$ stress -c 8 -i 4 -m 6 -t 20s
```

will:

- Fork off 8 CPU-intensive processes, each spinning on a `sqrt()` calculation.
- Fork off 4 I/O-intensive processes, each spinning on `sync()`.
- Fork off 6 memory-intensive processes, each spinning on `malloc()`, allocating 256 MB by default. The size can be changed as in `--vm-bytes 128M`.
- Run the stress test for 20 seconds.

After installing **stress**, you may want to start up your system's graphical system monitor, which you can find on your application menu, or run from the command line, which is probably **gnome-system-monitor** or **ksysguard**.

Now begin to put stress on the system. The exact numbers you use will depend on your system's resources, such as the number of CPUs and RAM size.

For example, doing

```
$ stress -m 4 -t 20s
```

puts only a memory stressor on the system.

Play with combinations of the switches and see how they impact each other. You may find the **stress** program useful to simulate various high load conditions.



## Chapter 23

# Process Monitoring



### Lab 23.1: Processes

1. Run **ps** with the options **-ef**. Then run it again with the options **aux**. Note the differences in the output.
2. Run **ps** so that only the process ID, priority, nice value, and the process command line are displayed.
3. Start a new **bash** session by typing **bash** at the command line. Start another **bash** session using the **nice** command but this time giving it a nice value of **10**.
4. Run **ps** as in step 2 to note the differences in priority and nice values. Note the process ID of the two **bash** sessions.
5. Change the nice value of one of the **bash** sessions to 15 using **renice**. Once again, observe the change in priority and nice values.
6. Run **top** and watch the output as it changes. Hit **q** to stop the program.

### Solution 23.1

1. 

```
$ ps -ef
$ ps aux
```
2. 

```
$ ps -o pid,pr,ni,cmd
```

| PID   | PR | NI | CMD                 |
|-------|----|----|---------------------|
| 2389  | 19 | 0  | bash                |
| 22079 | 19 | 0  | ps -o pid,pr,ni,cmd |

(Note: There should be no spaces between parameters.)
3. 

```
$ bash
$ nice -n 10 bash
$ ps -o pid,pr,ni,cmd
```

```

 2389 19 0 bash
22115 19 0 bash
22171 9 10 bash
22227 9 10 ps -o pid,pri,ni,cmd

4. $ renice 15 -p 22227
$ ps -o pid,pri,ni,cmd
 PID PRI NI CMD
 2389 19 0 bash
 22115 19 0 bash
 22171 4 15 bash
 22246 4 15 ps -o pid,pri,ni,cmd

5. $ top

```

### Lab 23.2: Monitoring Process States

1. Use **dd** to start a background process which reads from `/dev/urandom` and writes to `/dev/null`.
2. Check the process state. What should it be?
3. Bring the process to the foreground using the **fg** command. Then hit **Ctrl-Z**. What does this do? Look at the process state again, what is it?
4. Run the **jobs** program. What does it tell you?
5. Bring the job back to the foreground, then terminate it using **kill** from another window.

### Solution 23.2

```

1. $ dd if=/dev/urandom of=/dev/null &

2. $ ps -C dd -o pid,cmd,stat
25899 dd if=/dev/urandom of=/dev/ R

Should be S or R.

3. $ fg
$ ^Z
$ ps -C dd -o pid,cmd,stat
 PIDCMD STAT
25899 dd if=/dev/urandom of=/dev/ T

State should be T.

4. Type the jobs command. What does it tell you?

$ jobs

[1]+ Stopped dd if=/dev/urandom of=/dev/null

5. Bring the job back to the foreground, then kill it using the kill command from another window.

$ fg
$ kill 25899

```

## Chapter 24

# I/O Monitoring and Tuning



### Lab 24.1: bonnie++

**bonnie++** is a widely available benchmarking program that tests and measures the performance of drives and filesystems. It is descended from **bonnie**, an earlier implementation.

Results can be read from the terminal window or directed to a file, and also to a CSV format (comma-separated value). Companion programs, **bonnie++html** and **bonnie++text**, can be used convert to HTML and plain text output formats.

We recommend you read the **man** page for **bonnie++** before using as it has quite a few options regarding which tests to perform and how exhaustive and stressful they should be. A quick synopsis is obtained with:

```
$ bonnie++ -help

bonnie++: invalid option -- 'h'
usage:
bonnie++ [-d scratch-dir] [-c concurrency] [-s size(MiB)][:chunk-size(b)]
 [-n number-to-stat[:max-size[:min-size]][:num-directories[:chunk-size]]]
 [-m machine-name] [-r ram-size-in-MiB]
 [-x number-of-tests] [-u uid-to-use:gid-to-use] [-g gid-to-use]
 [-q] [-f] [-b] [-p processes | -y] [-z seed | -Z random-file]
 [-D]

Version: 1.96
```

A quick test can be obtained with a command like:

```
$ time sudo bonnie++ -n 0 -u 0 -r 100 -f -b -d /mnt
```

where:

- **-n 0** means don't perform the file creation tests.

- `-u 0` means run as root.
- `-r 100` means pretend you have 100 MB of RAM.
- `-f` means skip per character I/O tests.
- `-b` means do a **fsync** after every write, which forces flushing to disk rather than just writing to cache.
- `-d /mnt` just specifies the directory to place the temporary file created; make sure it has enough space, in this case 300 MB, available.

If you don't supply a figure for your memory size, the program will figure out how much the system has and will create a testing file 2-3 times as large. We are not doing that here because it takes much longer to get a feel for things.

On an **RHEL 7** system:

```
$ time sudo bonnie++ -n 0 -u 0 -r 100 -f -b -d /mnt

Using uid:0, gid:0.
Writing intelligently...done
Rewriting...done
Reading intelligently...done
start 'em...done...done...done...done...done...
Version 1.96 -----Sequential Output----- --Sequential Input-- --Random-
Concurrency 1 -Per Chr- --Block-- -Rewrite- -Per Chr- --Block-- --Seeks--
Machine Size K/sec %CP K/sec %CP K/sec %CP K/sec %CP K/sec %CP /sec %CP
q7 300M 99769 14 106000 12 ++++++++ 257.3 1
Latency 226us 237us 418us 624ms

1.96,1.96,q7,1,1415992158,300M,,,99769,14,106000,12,,,++++,+++257.3,1,,,,,,,,,,,,,226us,237us,,418us,624ms,,,,,
```

On an **Ubuntu 14.04** system, running as a virtual machine under hypervisor on the same physical machine:

```
$ time sudo bonnie++ -n 0 -u 0 -r 100 -f -b -d /mnt

Using uid:0, gid:0.
Writing intelligently...done
Rewriting...done
Reading intelligently...done
start 'em...done...done...done...done...done...
Version 1.97 -----Sequential Output----- --Sequential Input-- --Random-
Concurrency 1 -Per Chr- --Block-- -Rewrite- -Per Chr- --Block-- --Seeks--
Machine Size K/sec %CP K/sec %CP K/sec %CP K/sec %CP K/sec %CP /sec %CP
ubuntu 300M 70000 61 43274 31 470061 96 2554 91
Latency 306ms 201ms 9276us 770ms

1.97,1.97,ubuntu,1,1415983257,300M,,,70000,61,43274,31,,,470061,96,2554,91,,,,,,,,,,,,,306ms,201ms,,9276us,770ms,,,
```

You can clearly see the drop in performance.

Assuming you have saved the previous outputs as a file called `bonnie++.out`, you can convert the output to html:

```
$ bon_csv2html < bonnie++.out > bonnie++.html
```

or to plain text with:

```
$ bon_csv2txt < bonnie++.out > bonnie++.txt
```

After reading the document ation, try longer and larger, more ambit ious tests. Try some of the tests we turned off. If your system is behaving well, save the results for future benchmarking comparisons when the system is sick.

## Lab 24.2: fs mark

The **fs\_mark** benchmark gives a low level bashing to file systems, using heavily asynchronous I/O across multiple directories and drives. It's a rather old program written by Ric Wheeler that has stood the test of time.

It can be downloaded from <http://sourceforge.net/projects/fsmark/> Once you have obtained the tarball, you can unpack it and compile it with:

```
$ tar zxvf fs_mark-3.3.tgz
$ cd fs_mark
$ make
```

Read the **README** file as we are only going to touch the surface.

If the compile fails with an error like:

```
$ make
....
/usr/bin/ld: cannot find -lc
```

it is because you haven't installed the **static** version of **glibc**. You can do this on **Red Hat**-based systems by doing:

```
$ sudo yum install glibc-static
```

and on **SUSE**-related systems with:

```
$ sudo zypper install glibc-devel-static
```

On **Debian**-based systems the relevant static library is installed along with the shared one so no additional package needs to be sought.

For a test we are going to create 1000 files, each 10 KB in size, and after each write we'll perform an **fsync** to flush out to disk. This can be done in the **/tmp** directory with the command:

```
$ fs_mark -d /tmp -n 1000 -s 10240
```

While this is running, gather extended **iostat** statistics with:

```
$ iostat -x -d /dev/sda 2 20
```

in another terminal window.

The numbers you should surely note are the number of files per second reported by **fs\_mark** and the percentage of CPU time utilized reported by **iostat**. If this is approaching 100 percent, you are I/O-bound.

Depending on what kind of filesystem you are using you may be able to get improved results by changing the **mount** options. For example, for **ext3** or **ext4** you can try:

```
$ mount -o remount,barrier=1 /tmp
```

or for **ext4** you can try:

```
$ mount -o remount,journal_async_commit /tmp
```

See how your results change.

Note that these options may cause problems if you have a power failure, or other ungraceful system shutdown; i.e., there is likely to be a trade-off between stability and speed.

Documentation about some of the **mount** options can be found with the kernel source under [Documentation/filesystems](#) and the **man** page for **mount**.



## Chapter 25

# I/O Scheduling



### Lab 25.1: Comparing I/O Schedulers

We provide a script which is to be used to compare I/O schedulers:

```
#!/bin/bash

#/*
* The code herein is: Copyright the Linux Foundation, 2014
* Author J. Cooperstein
*
* This Copyright is retained for the purpose of protecting free
* redistribution of source.
*
* This code is distributed under Version 2 of the GNU General Public
* License, which you should have received with the source.
*
*/

NMAX=8
NMEGS=100
[[-n $1]] && NMAX=$1
[[-n $2]] && NMEGS=$2

echo Doing: $NMAX parallel read/writes on: $NMEGS MB size files

TIMEFORMAT="%R %U %S"

#####
simple test of parallel reads
do_read_test(){
 for n in $(seq 1 $NMAX) ; do
 cat file$n > /dev/null &
 done
```

```

wait for previous jobs to finish
wait
}

simple test of parallel writes
do_write_test(){
 for n in $(seq 1 $NMAX) ; do
 [[-f fileout$n]] && rm -f fileout$n
 (cp file1 fileout$n && sync) &
 done
wait for previous jobs to finish
wait
}

create some files for reading, ok if they are the same
create_input_files(){
 [[-f file1]] || dd if=/dev/urandom of=file1 bs=1M count=$NMEGS
 for n in $(seq 1 $NMAX) ; do
 [[-f file$n]] || cp file1 file$n
 done
}

echo -e "\ncreating as needed random input files"
create_input_files

#####
begin the actual work

do parallel read test
echo -e "\ndoing timings of parallel reads\n"
echo -e " REAL USER SYS\n"
for iosched in noop deadline cfq ; do
 echo testing IOSCHED = $iosched
 echo $iosched > /sys/block/sda/queue/scheduler
 cat /sys/block/sda/queue/scheduler
echo -e "\nclearing the memory caches\n"
echo 3 > /proc/sys/vm/drop_caches
 time do_read_test
done
#####
do parallel write test
echo -e "\ndoing timings of parallel writes\n"
echo -e " REAL USER SYS\n"
for iosched in noop deadline cfq ; do
 echo testing IOSCHED = $iosched
 echo $iosched > /sys/block/sda/queue/scheduler
 cat /sys/block/sda/queue/scheduler
 time do_write_test
done
#####

```

If you are taking the online self-paced version of this course, the script is available for download from your **Lab** screen.

Remember to make it executable by doing: by doing:

```
$ chmod +x ioscript.sh
```

The following explains how the script was written and how to use it.

The script should:

- Cycle through the available I/O schedulers on a hard disk while doing a configurable number of parallel reads and writes of files of a configurable size.



- Test reads and writes as separate steps.
- When testing reads make sure you're actually reading from disk and not from cached pages of memory; you can flush out the cache by doing:  

```
$ echo 3 > /proc/sys/vm/drop_caches
```

before doing the reads. You can **cat** into `/dev/null` to avoid writing to disk.
- Make sure all reads are complete before obtaining timing information; this can be done by issuing a **wait** command under the shell.
- Test writes by simply copying a file (which will be in cached memory after the first read) multiple times simultaneously. To make sure you wait for all writes to complete before you get timing information you can issue a **sync** call.

The provided script takes two arguments. The first is the number of simultaneous reads and writes to perform. The second is the size (in MB) of each file.

This script must be run as root as it echoes values into the `/proc` and `/sys` directory trees.

Compare the results you obtain using different I/O schedulers.

For additional exploring you might try changing some of the tuneable parameters and see how results vary.



## Chapter 26

# Memory: Monitoring Usage and Tuning



### Lab 26.1: Invoking the OOM Killer

Examine what swap partitions and files are present on your system by examining `/proc/swaps`.

Turn off all swap with the command

```
$ sudo /sbin/swapoff -a
```

Make sure you turn it back on later, when we are done, with

```
$ sudo/sbin/swapon -a
```

Now we are going to put the system under increasing memory pressure. One way to do this is to exploit the `stress` program we installed earlier, running it with arguments such as:

```
$ stress -m 8 -t 10s
```

which would keep 2 GB busy for 10 seconds.

You should see the **OOM** (Out of Memory) killer swoop in and try to kill processes in a struggle to stay alive. You can see what is going on by running `dmesg` or monitoring `/var/log/messages` or `/var/log/syslog`, or through graphical interfaces that expose the system logs.

Who gets clobbered first?



## Chapter 27

# Package Management Systems

There are no lab exercises in this chapter. It just sets the stage for the following sections on particular types of package management, each of which has labs.



## Chapter 28

# RPM



### Lab 28.1: Using RPM

Here we will just do a number of simple operations for querying and verifying **rpm** packages.

This lab will work equally well on **Red Hat** and **SUSE**-based systems.

1. Find out what package the file `/etc/logrotate.conf` belongs to.
2. List information about the package including all the files it contains.
3. Verify the package installation.
4. Try to remove the package.

### Solution 28.1

```
1. $ rpm -qf /etc/logrotate.conf
logrotate-3.8.6-4.el7.x86_64
```

```
2. $ rpm -qil logrotate
...
```

Note a fancier form that combines these two steps would be:

```
$ rpm -qil $(rpm -qf /etc/logrotate.conf)
```

```
3. $ rpm -V logrotate
..?..... /etc/cron.daily/logrotate
S.5....T. c /etc/logrotate.conf
```

4. On **RHEL 7**:

```
$ sudo rpm -e logrotate
error: Failed dependencies:
 logrotate is needed by (installed) vsftpd-3.0.2-9.el7.x86_64
 logrotate >= 3.5.2 is needed by (installed) rsyslog-7.4.7-7.el7_0.x86_64
```

On **openSUSE 13.1**:

```
$ sudo rpm -e logrotate
error: Failed dependencies:
 logrotate is needed by (installed) xdm-1.1.10-24.2.1.x86_64
 logrotate is needed by (installed) syslog-service-2.0-772.1.2.noarch
 logrotate is needed by (installed) wpa_supplicant-2.0-3.4.1.x86_64
 logrotate is needed by (installed) mcelog-1.0pre3.6e4e2a000124-19.4.1.x86_64
 logrotate is needed by (installed) apache2-2.4.6-6.27.1.x86_64
 logrotate is needed by (installed) net-snmp-5.7.2-9.8.1.x86_64
 logrotate is needed by (installed) kdm-4.11.12-119.1.x86_64
```

Note that the exact package dependency tree depends on both the distribution and choice of installed software.

## Lab 28.2: Rebuilding the RPM Database

There are conditions under which the **RPM** database stored in `/var/lib/rpm` can be corrupted. In this exercise we will construct a new one and verify its integrity.

This lab will work equally well on **Red Hat** and **SUSE**-based systems.

1. Backup the contents of `/var/lib/rpm` as the rebuild process will overwrite the contents. If you neglect to do this and something goes wrong you are in serious trouble.
2. Rebuild the data base.
3. Compare the new contents of the directory with the backed up contents; don't examine the actual file contents as they are binary data, but note the number and names of the files.
4. Get a listing of all **rpms** on the system. You may want to compare this list with one generated before you actually do the rebuild procedure. If the query command worked, your new database files should be fine.
5. Compare again the two directory contents. Do they have the same files now?
6. You could delete the backup (probably about 100 MB in size) but you may want to keep it around for a while to make sure your system is behaving properly before trashing it.

You may want to look at <http://www.rpm.org/wiki/Docs/RpmRecovery> for a more complete examination of steps you can take to verify and/or recover the database integrity.

## Solution 28.2

- ```
$ cd /var/lib
1. $ sudo cp -a rpm rpm_BACKUP
2. $ sudo rpm --rebuilddb
3. $ ls -l rpm rpm_BACKUP
4. $ rpm -qa | tee /tmp/rpm-qa.output
```


5. `$ ls -l rpm rpm_BACKUP`

6. Probably you should not do this until you are sure the system is fine!

`$ sudo rm -rf rpm_BACKUP`

Chapter 29

DPKG



Lab 29.1: Using dpkg

Here we will just do a number of simple operations for querying and verifying **Debian** packages.

1. Find out what package the file `/etc/logrotate.conf` belongs to.
2. List information about the package including all the files it contains.
3. Verify the package installation.
4. Try to remove the package.

Solution 29.1

1.

```
$ dpkg -S /etc/logrotate.conf
```



```
logrotate: /etc/logrotate.conf
```
2.

```
$ dpkg -L logrotate
```



```
...
```
3.

```
$ dpkg -V logrotate
```
4.

```
$ sudo dpkg -r logrotate
```



```
dpkg: dependency problems prevent removal of logrotate:
  libvirt-bin depends on logrotate.
  ubuntu-standard depends on logrotate.

dpkg: error processing package logrotate (--remove):
```

```
dependency problems - not removing  
Errors were encountered while processing:  
logrotate
```

Chapter 30

yum



Lab 30.1: Basic YUM Commands

1. Check to see if there are any available updates for your system.
2. Update a particular package.
3. List all installed kernel-related packages, and list all installed or available ones.
4. Install the **httpd-devel** package, or anything else you might not have installed yet. Doing a simple:

```
$ sudo yum list
```

will let you see a complete list; you may want to give a wildcard argument to narrow the list.

Solution 30.1

1.

```
$ sudo yum update
$ sudo yum check-update
$ sudo yum list updates
```

Only the first form will try to do the installations.
2.

```
$ sudo yum update bash
```
3.

```
$ sudo yum list installed "kernel*"
$ sudo yum list "kernel*"
```
4.

```
$ sudo yum install httpd-devel
```

Lab 30.2: Using yum to Find Information About a Package

Using **yum** (and not **rpm** directly), find:

1. All packages that contain a reference to **bash** in their name or description.
2. Installed and available **bash** packages.
3. The package information for **bash**.
4. The dependencies for the **bash** package.

Try the commands you used above both as **root** and as a regular user. Do you notice any difference?

Solution 30.2

Note: on **RHEL 7** you may get some permission errors if you don't use **sudo** with the following commands, even though we are just getting information.

1. `$ sudo yum search bash`
2. `$ sudo yum list bash`
3. `$ sudo yum info bash`
4. `$ sudo yum deplist bash`

All the commands above should work for both regular users and the **root** user.

Lab 30.3: Managing Groups of Packages with yum

Note: on **RHEL 7** you may get some permission errors if you don't use **sudo** with some of the following commands, even when we are just getting information.

yum provides the ability to manage groups of packages.

1. Use the following command to list all package groups available on your system:
`$ yum grouplist`
2. Identify the **Backup Client** group and generate the information about this group using the command
`$ yum groupinfo "Backup Client"`
3. Install using:
`$ sudo yum groupinstall "Backup Client"`
4. Identify a package group that's currently installed on your system and that you don't need. Remove it using **yum groupremove** as in:
`$ sudo yum groupremove "Backup Client"`

Note you will be prompted to confirm removal so you can safely type the command to see how it works.

You may find that the **groupremove** does **not** remove everything that was installed; whether this is a bug or a feature can be discussed.

Lab 30.4: Adding a New yum Repository

According to its authors (at <http://www.webmin.com/index.htm>):

“**Webmin** is a web-based interface for system administration for Unix. Using any modern web browser , you can setup user accounts, Apache, DNS, file sharing and much more. Webmin removes the need to manually edit Unix configuration files like `/etc/passwd`, and lets you manage a system from the console or remotely.”

We are going to create a repository for installation and upgrade. While we could simply go to the download page and get the current **rpm**, that would not automatically give us any upgrades.

1. Create a new repository file called `webmin.repo` in the `/etc/yum.repos.d` directory. It should contain the following:

```
[Webmin]
name=Webmin Distribution Neutral
baseurl=http://download.webmin.com/download/yum
mirrorlist=http://download.webmin.com/download/yum/mirrorlist
enabled=1
gpgcheck=0
```

(Note you can also cut and paste the contents from <http://www.webmin.com/download.html>)

2. Install the webmin package.

```
$ sudo yum install webmin
```


Chapter 31

zypper



Lab 31.1: Basic zypper Commands

1. Check to see if there are any available updates for your system.
2. Update a particular package.
3. List all repositories the system is aware of, enabled or not.
4. List all installed kernel-related packages, and list all installed or available ones.
5. Install the **apache2-devel** package, or anything else you might not have installed yet. (Note **httpd** is **apache2** on **SUSE** systems.) Doing a simple:

```
$ sudo zypper search
```

will let you see a complete list; you may want to give a wildcard argument to narrow the list.

Solution 31.1

1. `$ zypper list-updates`
2. `$ sudo zypper update bash`
3. `$ zypper repos`
4. `$ zypper search -i kernel`
`$ zypper search kernel`
5. `$ sudo zypper install apache2-devel`

Lab 31.2: Using zypper to Find Information About a Package

Using **zypper** (and not **rpm** directly), find:

1. All packages that contain a reference to **bash** in their name or description.
2. Installed and available **bash** packages.
3. The package information for **bash**.
4. The dependencies for the **bash** package.

Try the commands you used above both as **root** and as a regular user. Do you notice any difference?

Solution 31.2

1. `$ zypper search -d bash`

Without the `-d` option only packages with **bash** in their actual name are reported. You may have to do `zypper info` on the package to see where **bash** is mentioned.

2. `$ zypper search bash`

3. `$ zypper info bash`

4. `$ zypper info--requires bash`

will give a list of files **bash** requires. Perhaps the easiest way to see what depends on having **bash** installed is to do

```
$ sudo zypper remove --dry-run bash
```

For this exercise **bash** is a bad choice since it is so integral to the system; you really can't remove it anyway.

Chapter 32

APT



Lab 32.1: Basic APT Commands

1. Check to see if there are any available updates for your system.
2. Update a particular package.
3. List all installed kernel-related packages, and list all installed or available ones.
4. Install the **apache2-dev** package, or anything else you might not have installed yet. Doing a simple:

```
$ apt-cache pkgnames
```

will let you see a complete list; you may want to give a wildcard argument to narrow the list.

Solution 32.1

1. First synchronize the package index files with remote repositories:

```
$ sudo apt-get update
```

To actually upgrade:

```
$ sudo apt-get upgrade  
$ sudo apt-get -u upgrade
```

(You can also use **dist-upgrade** as discussed earlier.) Only the first form will try to do the installations.

2.

```
$ sudo apt-get upgrade bash
```
3.

```
$ apt-cache search "kernel"  
$ apt-cache search -n "kernel"  
$ apt-cache pkgnames "kernel"
```

The second and third forms only find packages that have `kernel` in their name.

```
$ dpkg --get-selections "*kernel*"
```

to get only installed packages. Note that on **Debian**-based systems you probably should use `linux` not `kernel` for kernel-related packages as they don't usually have `kernel` in their name.

4. `$ sudo apt-get install apache2-d ev`

Lab 32.2: Using APT to Find Information About a Package

Using `apt-cache` and `apt-get` (and not `dpkg`), find:

1. All packages that contain a reference to `bash` in their name or description.
2. Installed and available `bash` packages.
3. The package information for `bash`.
4. The dependencies for the `bash` package.

Try the commands you used above both as `root` and as a regular user. Do you notice any difference?

Solution 32.2

1. `$ apt-cache search bash`
2. `$ apt-cache search -n bash`
3. `$ apt-cache show bash`
4. `$ apt-cache depends bash`
`$ apt-cache rdepends bash`

Lab 32.3: Managing Groups of Packages with APT

APT provides the ability to manage groups of packages, similar to the way **yum** does it, through the use of **metapackages**. These can be thought of as **virtual packages**, that collect related packages that must be installed and removed as a group.

To get a list of available **metapackages**

```
$ apt-cache search metapackage
```

```
bacula - network backup service - metapackage
bacula-client = network backup service - client metapackage
cloud-utils - metapackage for installation of upstream cloud-utils source
compiz - OpenGL window and compositing manager
emacs - GNU Emacs editor (metapackage)
....
```

You can then easily install them like regular single packages, as in:

```
$ sudo apt-get install bacula-client

Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  bacula-common bacula-console bacula-fd bacula-traymonitor
Suggested packages:
  bacula-doc kde gnome-desktop-environment
The following NEW packages will be installed:
  bacula-client bacula-common bacula-console bacula-fd bacula-traymonitor
0 upgraded, 5 newly installed, 0 to remove and 0 not upgraded.
Need to get 742 kB of archives.
After this operation, 1,965 kB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Select an uninstalled metapackage and then remove it.

Chapter 33

User Account Management



Lab 33.1: Working with User Accounts

1. Examine `/etc/passwd` and `/etc/shadow`, comparing the fields in each file, especially for the normal user account. What is the same and what is different?
2. Create a `user1` account using `useradd`.
3. Login as `user1` using `ssh`. You can just do this with:

```
$ ssh user1@localhost
```

It should fail because you need a password for `user1`; it was never established.
4. Set the password for `user1` to `user1pw` and then try to login again as `user1`.
5. Look at the new records which were created in the `/etc/passwd`, `/etc/group` and the `/etc/shadow` files.
6. Look at the `/etc/default/useradd` file and see what the current defaults are set to. Also look at the `/etc/login.defs` file.
7. Create a user account for `user2` which will use the **Korn** shell (`ksh`) as its default shell. (if you don't have `/bin/ksh` install it or use the **C** shell at `/bin/csh`.) Set the password to `user2pw`.
8. Look at `/etc/shadow`. What is the current expiration date for the `user1` account?
9. Use `chage` to set the account expiration date of `user1` to December 1, 2013.

Look at `/etc/shadow` to see what the new expiration date is.
10. Use `usermod` to lock the `user1` account.

Look at `/etc/shadow` and see what has changed about `user1`'s password. Reset the password to `userp1` on the account to complete this exercise.

Solution 33.1

```
1. $ sudo grep student /etc/passwd /etc/shadow
/etc/passwd:student:x:1000:100:LF Student:/home/student:/bin/bash
/etc/shadow:student:$6$jtoFVPICHHba$iGFFU8ctrtrt0GoistJ4/30DrNLi1FS66qnm0VbS6Mvm
lcKI08SgbzT5.Ic0Ho5j/S0dCagZmF2RgzTzvZLb11H0:16028:0:99999:7:::
```

(You can use any normal user name in the place of `student`.) About the only thing that matches is the user name field.

```
2. $ sudo useradd user1
```

```
3. $ ssh user1@localhost
user1@localhost's password:
```

Note you may have to first start up the `sshd` service as in:

```
$ sudo service sshd restart
```

or

```
$ sudo systemctl restart sshd.service
```

```
4. $ sudo passwd user1
Changing password for user user1.
New password:
```

```
5. $ sudo grep user1 /etc/passwd /etc/shadow
/etc/passwd:user1:x:1001:100:./home/user1:/bin/bash
/etc/shadow:user1:$6$0BE1mPMw$CIc7urbQ9ZSnyiniV0eJxKqLFu8fz4whfEexVem2
TFpucuWRN1CCHZ19XGhj4qVujslRIS.P4aCXd/y1U4utv.:16372:0:99999:7:::
```

6. On either **RHEL 7** or **openSUSE 13.1** systems for example:

```
$ cat /etc/default/useradd
# useradd defaults file
GROUP=100
HOME=/home
INACTIVE=-1
EXPIRE=
SHELL=/bin/bash
SKEL=/etc/skel
CREATE_MAIL_SPOOL=yes
$ cat /etc/login.defs
....
```

We don't reproduce the second file as it is rather longer, but examine it on your system.

```
7. $ sudo useradd -s /bin/ksh user2
$ sudo passwd user2
```

```
Changing password for user user2.
New password:
```

```
8. $ sudo grep user1 /etc/shadow
user1:$6$0BE1mPMw$CIc7urbQ9ZSnyiniV0eJxKqLFu8fz4whfEexVem2TFpucuWRN1CCHZ
19XGhj4qVujslRIS.P4aCXd/y1U4utv.:16372:0:99999:7:::
```

There should be no expiration date.

9.

```
$ sudo chage -E 2013-12-1 user1
$ sudo sudo grep user1 /etc/shadow
user1:$6$0BE1mPMw$CIc7urbQ9ZSnyiniV0eJxKqLFu8fz4whfEexVem2TFpucuWRN1CCHZ
19XGhj4qVujs1RIS.P4aCXd/y1U4utv.:16372:0:99999:7::16040:
```
10.

```
$ sudo usermod -L user1

$ sudo passwd user1
```


Chapter 34

Group Management



Lab 34.1: Working with Groups

1. Create two new user accounts (**rocky** and **bullwinkle** in the below) and make sure they have home directories.
2. Create two new groups, **friends** and **bosses** (with a GID of 490). Look at the `/etc/group` file. See what GID was given to each new group.
3. Add **rocky** to both new groups.
Add **bullwinkle** to group **friends**.
Look in the `/etc/group` file to see how it changed.
4. Login as **rocky**. Create a directory called **somedir** and set the group ownership to **bosses**. (Using **chgroup** which will be discussed in the next session.)
(You will probably need to add execute privileges for all on **rocky**'s home directory.)
5. Login as **bullwinkle** and try to create a file in `/home/rocky/somedir` called **somefile** using the **touch** command.
Can you do this? No, because of the group ownership and the `chmod a+x` on the directory.
6. Add **bullwinkle** to the **bosses** group and try again `.`. Note you will have to logout and log back in again for the new group membership to be effective. do the following:

Solution 34.1

1.

```
$ sudo useradd -m rocky
$ sudo useradd -m bullwinkle
$ sudo passwd rocky

Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
```

```

$ sudo passwd bullwinkle
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully

$ ls -l /home
total 12
drwxr-xr-x  2 bullwinkle bullwinkle 4096 Oct 30 09:39 bullwinkle
drwxr-xr-x  2 rocky      rocky      4096 Oct 30 09:39 rocky
drwxr-xr-x 20 student    student    4096 Oct 30 09:18 student

```

2. \$ sudo groupadd friends
 \$ sudo groupadd -g 490 bosses
 \$ grep -e friends -e bosses /etc/group

```

friends:x:1003:
bosses:x:490:

```

3. \$ sudo usermod -G friends,bosses rocky
 \$ sudo usermod -G friends bullwinkle

```

$ grep -e rocky -e bullwinkle /etc/group
rocky:x:1001:
bullwinkle:x:1002:
friends:x:1003:rocky,bullwinkle
bosses:x:490:rocky

$ groups rocky bullwinkle
rocky : rocky friends bosses
bullwinkle : bullwinkle friends

```

4. \$ ssh rocky@localhost
 \$ cd ~
 \$ mkdir somedir
 \$ chgrp bosses somedir
 \$ ls -l

```

total 16
-rw-r--r-- 1 rocky rocky 8980 Oct  4 2013 examples.desktop
drwxrwxr-x 2 rocky bosses 4096 Oct 30 09:53 somedir

$ chmod a+x .

```

5. \$ ssh bullwinkle@localhost
 \$ touch /home/rocky/somedir/somefile

```

touch: cannot touch /home/rocky/somedir/somefile: Permission denied
$ exit

```

6. \$ sudo usermod -a -G bosses bullwinkle
 \$ ssh bullwinkle@localhost
 \$ touch /home/rocky/somedir/somefile
 \$ ls -al /home/rocky/somedir

(note ownership of files)

Chapter 35

File Permissions and Ownership



Lab 35.1: Using chmod

One can use either the octal digit or symbolic methods for specifying permissions when using **chmod**. Let's elaborate some more on the symbolic method.

It is possible to either give permissions directly, or add or subtract permissions. The syntax is pretty obvious. Try the following examples:

```
$ chmod u=r,g=w,o=x afile
$ chmod u+=w,g=-w,o=+rw afile
$ chmod ug=rwx,o=-rw afile
```

After each step do:

```
$ ls -l afile
```

to see how the permissions took, and try some variations.

Lab 35.2: umask

Create an empty file with:

```
$ touch afile
$ ls -l afile

-rw-rw-r-- 1 coop coop 0 Jul 26 12:43 afile
```

which shows it is created by default with both read and write permissions for owner and group, but only read for world.

In fact, at the operating system level the default permissions given when creating a file or directory are actually read/write for owner, group **and** world (0666); the default values have actually been modified by the current **umask**.

If you just type **umask** you get the current value:

```
$ umask
```

```
0002
```

which is the most conventional value set by system administrators for users. This value is combined with the file creation permissions to get the actual result; i.e.,

```
0666 & ~002 = 0664; i.e., rw-rw-r--
```

Try modifying the **umask** and creating new files and see the resulting permissions, as in:

```
$ umask 0022
$ touch afile2
$ umask 0666
$ touch afile3
$ ls -l afile*
```

Chapter 36

Pluggable Authentication Modules (PAM)



Lab 36.1: PAM Configuration

One of the more common PAM configurations is to deny login access after a certain number of failed attempts. This is done with the `pam_tally2` module. In this exercise we are going to deny login through `ssh` after three failed login attempts.

1. Edit `/etc/pam.d/ssh` and configure it to deny login after three failed attempts. Hint: add the following two lines to the file:

```
auth required pam_tally2.so deny=3 onerr=fail
account required pam_tally2.so
```
2. Try to login three times as a particular user (who has an account) while mistyping the password.
3. Try to login as the same user with the correct password.
4. Check to see how many failed logins there are for the user.
5. Reset the failed login counter.
6. Check again to see how many failed logins there are.
7. Try to login again with the correct password.

Solution 36.1

1. Add the following two lines to `/etc/pam.d/ssh`:

```
auth required pam_tally2.so deny=3 onerr=fail
account required pam_tally2.so
```

```
2. $ ssh student@localhost
Password:
Password:
Password:
Permission denied (publickey,keyboard-interactive).

3. $ ssh student@localhost
Password:
Account locked due to 3 failed logins

4. $ sudo pam_tally2
Login      Failures Latest failure    From
student          3   11/01/14 20:41:12 localhost

5. $ sudo pam_tally2 -u student -r
Login      Failures Latest failure    From
student          3   11/01/14 20:41:12 localhost

6. $ sudo pam_tally2 -u student -r
Login      Failures Latest failure    From
student          0

7. $ ssh student@localhost
Password:
Last failed login: Sat Nov  1 20:41:14 CDT 2014 from localhost on ssh:notty
There were 6 failed login attempts since the last successful login.
Last login: Sat Nov  1 20:28:38 2014 from localhost
Have a lot of fun...
```


Chapter 37

Backup and Recovery Methods



Lab 37.1: Using tar for Backup

1. Create a directory called **backup** and in it place a compressed **tar** archive of all the files under **/usr/include**, with the highest level directory being **include**. You can use any compression method (**gzip**, **bzip2** or **xzip**).
2. List the files in the archive.
3. Create a directory called **restore** and unpack and decompress the archive.
4. Compare the contents with the original directory the archive was made from.

Solution 37.1

1.

```
$ cd backup
$ cd /usr ; tar zcvf include.tar.gz include
```

or

```
$ tar -C /usr -zcf include.tar.gz include
$ tar -C /usr -jcf include.tar.bz2 include
$ tar -C /usr -Jcf include.tar.xz include
```

Notice the efficacy of the compression between the three methods:

```
$ du -sh /usr/include
55M      /usr/include
```

2.

```
$ ls -lh include.tar.*
-rw-rw-r-- 1 coop coop 5.3M Nov  3 14:44 include.tar.bz2
-rw-rw-r-- 1 coop coop 6.8M Nov  3 14:44 include.tar.gz
-rw-rw-r-- 1 coop coop 4.7M Nov  3 14:46 include.tar.xz
```

3. `$ tar tvf include.tar.xz`

```

qdrwxr-xr-x root/root      0 2014-10-29 07:04 include/
-rw-r--r-- root/root    42780 2014-08-26 12:24 include/unistd.h
-rw-r--r-- root/root      957 2014-08-26 12:24 include/re_comp.h
-rw-r--r-- root/root   22096 2014-08-26 12:24 include/regex.h
-rw-r--r-- root/root     7154 2014-08-26 12:25 include/link.h
.....

```

Note it is not necessary to give the `j`, `J`, or `z` option when decompressing; `tar` is smart enough to figure out what is needed.

4. `$ cd .. ; mkdir restore ; cd restore`

```

$ tar xvf ../backup/include.tar.bz2
include/
include/unistd.h
include/re_comp.h
include/regex.h
include/link
.....
$ diff -qr include /usr/include

```

Lab 37.2: Using cpio for Backup

We are going to do essentially the same exercise now, but using `cpio` in place of `tar`. We'll repeat the slightly altered instructions for ease of use.

1. Create a directory called `backup` and in it place a compressed `cpio` archive of all the files under `/usr/include`, with the highest level directory being `include`. You can use any compression method (`gzip`, `bzip2` or `xzip`).
2. List the files in the archive.
3. Create a directory called `restore` and unpack and decompress the archive.
4. Compare the contents with the original directory the archive was made from.

Solution 37.2

1. `$ (cd /usr ; find include | cpio -c -o > /home/student/backup/include.cpio)`

```
82318 blocks
```

or to put it in a compressed form:

```
$ (cd /usr ; find include | cpio -c -o | gzip -c > /home/student/backup/include.cpio.gz)
```

```
82318 blocks
```

```
$ ls -lh include*
```

```

total 64M
-rw-rw-r-- 1 coop coop 41M Nov  3 15:26 include.cpio
-rw-rw-r-- 1 coop coop 6.7M Nov  3 15:28 include.cpio.gz
-rw-rw-r-- 1 coop coop 5.3M Nov  3 14:44 include.tar.bz2
-rw-rw-r-- 1 coop coop 6.8M Nov  3 14:44 include.tar.gz
-rw-rw-r-- 1 coop coop 4.7M Nov  3 14:46 include.tar.xz

```

2. `$ cpio -ivt < include.cpio`

```
drwxr-xr-x 86 root    root      0 Oct 29 07:04 include
-rw-r--r-- 1 root    root    42780 Aug 26 12:24 include/unistd.h
-rw-r--r-- 1 root    root     957 Aug 26 12:24 include/re_comp.h
-rw-r--r-- 1 root    root   22096 Aug 26 12:24 include/regex.h
.....
```

Note the redirection of input; the archive is not an argument. One could also do:

```
$ cd ../restore
$ cat ../backup/include.cpio | cpio -ivt
$ gunzip -c include.cpio.gz | cpio -ivt
```

3.

```
$ rm -rf include
$ cpio -id < ../backup/include.cpio
$ ls -lR include
```

or

```
$ cpio -idv < ../backup/include.cpio

$ diff -qr include /usr/include
```

Lab 37.3: Using rsync for Backup

1. Using **rsync**, we will again create a complete copy of `/usr/include` in your backup directory:

```
$ rm -rf include
$ rsync -av /usr/include .
sending incremental file list
include/
include/FlexLexer.h
include/_G_config.h
include/a.out.h
include/aio.h
.....
```

2. Let's run the comm and a second time and see if it does anything:

```
$ rsync -av /usr/include .
sending incremental file list

sent 127398 bytes  received 188 bytes  255172.00 bytes/sec
total size is 41239 979  speedup is 323.23
```

3. One confusing thing about **rsync** is you might have expected the right command to be:

```
$ rsync -av /usr/include include
sending incremental file list
...
```

However, if you do this, you'll find it actually creates a new directory, `include/include!`

4. To get rid of the extra files you can use the `--delete` option:

```
$ rsync -av --delete /usr/include .
sending incremental file list
include/
deleting include/include/xen/privcmd.h
deleting include/include/xen/evtchn.h
....
deleting include/include/FlexLexer.h
deleting include/include/

sent 127401 bytes  received 191 bytes  85061.33 bytes/sec
total size is 41239 979  speedup is 323.22
```

5. For another simple exercise, remove a subdirectory tree in your backup copy and then run **rsync** again with and without the **--dry-run** option:

```
$ rm -rf include/xen
$ rsync -av --delete --dry-run /usr/include .
sending incremental file list
include/
include/xen/
include/xen/evtchn.h
include/xen/privcmd.h

sent 127412 bytes  received 202 bytes  255228.00 bytes/sec
total size is 41239 979  speedup is 323.16 (DRY RUN)
$ rsync -av --delete /usr/include .
```

6. A simple script with a good set of options for using **rsync**:

```
#!/bin/sh
set -x

rsync --progress -avrxH -e "ssh -c blowfish" --delete $*
```

which will work on a local machine as well as over the network. Note the important **-x** option which stops **rsync** from crossing filesystem boundaries.

For more fun, if you have access to more than one computer, try doing these steps with source and destination on different machines.

Chapter 38

Network Addresses

There are no lab exercises in this chapter. It just sets the stage for the following section on network configuration, which has several labs.

Chapter 39

Network Devices and Configuration



Lab 39.1: Static Configuration of a Network Interface

Note: you may have to use a different network interface name than `eth0`. Of course you can do this exercise from a graphical interface but we will present a command line solution.

1. Show your current IP address, default route and **DNS** settings for `eth0`. Keep a copy of them for resetting later.
2. Bring down `eth0` and reconfigure to use a static address instead of **DCHP**, using the information you just recorded.
3. Bring the interface back up, and configure the nameserver resolver with the information that you noted before.
Verify your hostname and then **ping** it.
4. Make sure your configuration works after a reboot.

You will probably want to restore your configuration when you are done.

Solution 39.1

1.

```
$ ifconfig eth0
$ route -n
$ cp /etc/resolv.conf resolv.conf.keep
```
2.

```
$ sudo ifconfig eth0 down
```

Make sure the following is in `/etc/sysconfig/network-scripts/ifcfg-eth0` on **Red Hat**-based systems:

```

DEVICE=eth0
BOOTPROTO=static
ONBOOT=yes
IPADDR=noted from step 1
NETMASK=noted from step 1
GATEWAY=noted from step 1

```

On **SUSE**-based systems edit the file in `/etc/sysconfig/network` in the same way, and on **Debian**-based systems edit `/etc/networking/interfaces` to include:

```

iface eth0 inet static
    address noted from step 1
    netmask noted from step 1
    gateway noted from step 1

```

3.

```
$ sudo ifconfig eth0 up
```



```
$ sudo cp resolv.conf.keep /etc/resolv.conf
```

```
$ cat /etc/sysconfig/network
```

```
$ cat /etc/hosts
```

```
$ ping yourhostname
```
4.

```
$ sudo reboot
```



```
$ ping hostname
```

Lab 39.2: Adding a Static Hostname

In this exercise we will add entries to the local host database.

1. Open `/etc/hosts` and add an entry for `mysystem.mydomain` that will point to the IP address associated with your network card.
2. Add a second entry that will make all references to `ad.doubleclick.net` point to `127.0.0.1`.
3. As an optional exercise, download the host file from: <http://winhelp2002.mvps.org/hosts2.htm> or more directly from <http://winhelp2002.mvps.org/hosts.txt> and install it on your system. Do you notice any difference using your browser with and without the new host file in place?

Solution 39.2

1. As root do:


```
$ echo "192.168.1.180    mysystem.mydomain" >> /etc/hosts
```

```
$ ping mysystem.mydomain
```
2. As root do:


```
$ echo "127.0.0.1      ad.doubleclick.net" >> /etc/hosts
```

```
$ ping ad.doubleclick.net
```
3.

```
$ wget http://winhelp2002.mvps.org/hosts.txt
```

```
--2014-11-01 08:57:12--  http://winhelp2002.mvps.org/hosts.txt
```

```
Resolving winhelp2002.mvps.org (winhelp2002.mvps.org)... 216.155.126.40
```

```
Connecting to winhelp2002.mvps.org (winhelp2002.mvps.org)|216.155.126.40|:80... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 514744 (503K) [text/plain]
```

```
Saving to: hosts.txt
```



```
100%[=====>] 514,744    977KB/s   in 0.5s
2014-11-01 08:57:13 (977 KB/s) - hosts.txt saved [514744/514744]

As root do:

$ cat hosts.txt >> /etc/hosts
```

Lab 39.3: Adding a Network Interface Alias

1. Configure your system with a new network device alias name `eth0:0`, which uses a new IP address you will select. This address should be persistent.
Bring the device up and test it.

Solution 39.3

1.

```
$ cd /etc/sysconfig/network-scripts
$ cp ifcfg-eth0 ifcfg-eth0:0
```

Edit this file (as root) and make sure it has the lines:

```
DEVICE=eth0:0
BOOTPROTO=static
ONBOOT=yes
IPADDR=192.168.1.110
NETMASK=255.255.255.0
```

using whatever address you want. On **RHEL 7** you should use `NAME` instead of `DEVICE`.

To bring the device up you can use `ifconfig`, `ifup` or `ip`, but simply doing:

```
$ sudo service network restart
```

will also show the new alias is persistent. You can test with

```
$ sudo ping 192.168.1.110
```

using whatever address you chose.

Chapter 40

Firewalls



Lab 40.1: Installing firewalld

While most recent **Linux** distributions have the **firewalld** package (which includes the **firewall-cmd** multi-purpose utility) available, it might not be installed on your system.

First you should check to see if it is already installed, with

```
$ which firewalld firewall-cmd
/usr/sbin/firewalld
/usr/bin/firewall-cmd
```

If you fail to find the program, then you need to install with one of the following, depending on your distribution in the usual way:

```
$ sudo yum install firewalld
$ sudo zypper install firewall
$ sudo apt-get install firewalld
```

If this fails, the **firewalld** package is not available for your distribution. For example, this will be the case for the older **RHEL6/CentOS6** distributions. In this case you will have to install from source.

To do this, go to <https://fedorahosted.org/firewalld/> and you can get the **git** source repository, but it is much easier to download the most recent tarball (**firewalld-0.3.13.tar.bz2** as of this writing.)

Then you have to follow the common procedure for installing from source:

```
$ tar xvf firewalld-0.3.13.tar.bz2
$ cd firewalld-0.3.13
$ ./configure
$ make
$ sudo make install
```

Note this source also has an **uninstall** target:

```
$ sudo make uninstall
```

in case you have regrets.

You will have to deal with any inadequacies that come up in the `./configure` step, such as missing libraries etc. When you install from a packaging system, the distribution takes care of this for you, but from source it can be problematic. If you have run the **Linux Foundation's** `ready-for.sh` script on your system, you are unlikely to have problems.

Note: On **openSUSE 13.2**, even though the compile and install of `firewalld` will work, execution of `firewall-cmd` will still fail with a message about missing `python-slip`. Unfortunately, this package also doesn't exist in the **zypper** repositories, so you will have to download it from the same web site, <https://fedorahosted.org/firewalld/>, and then just do:

```
$ tar xvf /tmp/python-slip-0.6.1.tar.bz2
$ cd python-slip-0.6.1
$ make
$ sudo make install
```

substituting the actual name of the version you downloaded. Hopefully, the next edition of **openSUSE** will eliminate this need to compile from sources, as there have been requests to add `firewalld` properly to the available choices.

Lab 40.2: Examining firewall-cmd

We have only scratched the surface of how you can use the `firewalld` package. Almost everything is done by deploying `firewall-cmd` which is empowered to do a large variety of tasks, using options with very clear names.

To get a sense of this, there is really no substitute for just doing:

```
$ firewall-cmd --help

Usage: firewall-cmd [OPTIONS...]
....
Service Options
  --new-service=<service>
                        Add a new service [P only]
  --delete-service=<service>
                        Delete an existing service [P only]
....
```

which we will not reproduce here as it is 208 lines on a **RHEL 7** system.

For more detailed explanation of anything which piques your interest, do `man firewall-cmd` which explains things more deeply, and `man firewalld` which gives an overview, as well as a listing of other `man` pages that describe the various configuration files in `/etc`, and elucidate concepts such as **zones** and **services**.

Lab 40.3: Adding Services to a Zone

Add the **http** and **https** services to the **public zone** and verify that they are currently listed.

Solution 40.3

```
$ sudo firewall-cmd --zone=public --add-service=http
success
$ sudo firewall-cmd --zone=public --add-service=https
```

```
success
$ sudo firewall-cmd --list-services --zone=public
dhcpv6-client http https ssh
```

Note if you had run

```
$ sudo firewall-cmd --reload
$ sudo firewall-cmd --list-services --zone=public
dhcpv6-client ssh
```

after adding the new services, they would disappear from the list! This curious behavior is because we did not include the `--permanent` flag when adding the services, and the `--reload` option reloads the known persistent services only.

Lab 40.4: Using the firewall GUI

Each distribution has its own graphical interface for firewall administration. On **Red Hat**-based systems you can run **firewall-config**, on **Ubuntu** it is called **gufw**, and on **openSUSE** you can find it as part of **yast** on the graphical menu system.

We have concentrated on the command line approach simply because we want to be distribution-flexible. However, for most relatively simple firewall configuration tasks, you can probably do them efficiently with less memorization from the GUI.

Once you launch the firewall configuration GUI, do the previous exercise of adding **http** and **https** to the **public** zone, and verify that it has taken effect.

Make sure you take the time to understand the graphical interface.

Chapter 41

Basic Troubleshooting

There are no lab exercises for in this chapter. It just summarizes points discussed earlier when considering configuring and monitoring the system, and in addition, sets the stage for the following section on system rescue, which has several labs.

Chapter 42

System Rescue



Lab 42.1: Preparing to use Rescue/Recover Media

In the following exercises we are going to deliberately damage the system and then recover through the use of rescue media. Thus, it is obviously prudent to make sure you can indeed boot off the rescue media before you try anything more ambitious.

So first make sure you have rescue media, either a dedicated rescue/recovery image, or an install or Live image on either an optical disk or usb drive.

Boot off it and make sure you know how to force the system to boot off the rescue media (you are likely to have to fiddle with the **BIOS** settings), and when the system boots, choose rescue mode.

If you are using a virtual machine, the procedure is logically the same with two differences:

- Getting to the **BIOS** might be difficult depending on the hypervisor you use. Some of them require very rapid keystrokes, so read the documentation and make sure you know how to do it.
- You can use a physical optical disk or drive, making sure the virtual machine settings have it mounted, and if it is **USB** you may have some other hurdles to make sure the virtual machine can claim the physical device. It is usually easier to simply connect a `.iso` image file directly to the virtual machine.

If you are working with a virtual machine, obviously things are less dangerous, and if you are afraid of corrupting the system in an unfixable way, simply make a backup copy of the virtual machine image before you do these exercises, you can always replace the image with it later.

Do not do the following exercises unless you are sure you can boot your system off rescue/recovery media!

Lab 42.2: Recovering from a Corrupted GRUB Configuration

1. Edit your **GRUB** configuration file (`/boot/grub/grub.cfg`, `/boot/grub2/grub.cfg` or `/boot/grub/grub.conf`), and modify the **kernel** line by removing the first character of the value in the field named `UUID`. Take note of which

- character you removed, you will replace it in rescue mode. (If your root filesystem is identified by either label or hard disk device node, make an analogous simple change.) Keep a backup copy of the original.
2. Reboot the machine. The system will fail to boot, saying something like **No root device** was found. You will also see that a **panic** occurred.
 3. Insert into your machine the **installation** or **Live DVD** or **CD** or **USB** drive (or network boot media) if you have access to a functioning installation server). Reboot again. When the boot menu appears, choose to enter rescue mode.
 4. As an alternative, you can try selecting a rescue image from the **GRUB** menu; most distributions offer this. You'll get the same experience as using rescue media, but it will not always work. For example, if the root filesystem is damaged it will be impossible to do anything.
 5. In rescue mode, agree when asked to search for filesystems. If prompted, open a shell, and explore the rescue system by running utilities such as **mount** and **ps**.
 6. Repair your broken system by fixing your **GRUB** configuration file, either by editing it or restoring from a backup copy.
 7. Type **exit** to return to the installer, remove the boot media, and follow the instructions on how to reboot. Reboot your machine. It should come up normally.

Lab 42.3: Recovering from Password Failure

1. As root (not with **sudo**), change the root password. We will pretend we don't know what the new password is.
2. Logout and try to login again as root using the old password. Obviously you will fail.
3. Boot using the rescue media, and select **Rescue** when given the option. Let it mount filesystems and then go to a command line shell.
4. Go into your **chroot**-ed environment (so you have normal access to your systems):


```
$ chroot /mnt/sysimage
```

 and reset the root password back to its original value.
5. Exit, remove the rescue media, and reboot, you should be able to login normally now.

Lab 42.4: Recovering from Partition Table Corruption

1. Login as root and save your **MBR**:

```
$ dd if=/dev/sda of=/root/mbrsave bs=446 count=1
1+0 records in
1+0 records out
446 bytes (446 B) copied, 0.00976759 s, 45.7 kB/s
```

Be careful: make sure you issue the exact command above and that the file saved has the right length:

```
$ sudo ls -l /root/mbrsave
-rw-r--r-- 1 root root 446 Nov 12 07:54 mbrsave
```

2. Now we are going to obliterate the **MBR** with:

```
$ dd if=/dev/zero of=/dev/sda bs=446 count=1
1+0 records in
1+0 records out
446 bytes (446 B) copied, 0.000124091 s, 3.6 MB/s
```

3. Reboot the system; it should fail.

4. Reboot into the rescue environment and restore the MBR:

```
$ dd if=/mnt/sysimage/root/mbrsave of=/dev/sda bs=446 count=1
```

5. Exit from the rescue environment and reboot. The system should boot properly now.

Lab 42.5: Recovering Using the Install Image

1. This exercise has been specifically written for **Red Hat**-based systems. You should be able to easily construct the appropriate substitutions for other distribution families.

Remove the **zsh** package (if it is installed!):

```
$ yum remove zsh
```

or

```
$ rpm -e zsh
```

Note we have chosen a package that generally has no dependencies to simplify matters. If you choose something that does, you will have to watch your step in the below so that anything else you remove you reinstall as needed as well.

2. Boot into the rescue environment.
3. Re-install (or install) **zsh** from within the rescue environment. First, mount the install media at `/mnt/source`:

```
$ mount /dev/cdrom /mnt/source
```

Then reinstall the package:

```
$ rpm -ivh --force --root /mnt/sysimage /mnt/source/Packages/zsh*.rpm
```

The `--force` option tells **rpm** to use the source directory in determining dependency information etc. Note that if the install image is much older than your system which has had many updates the whole procedure might collapse!

4. Exit and reboot.
5. Check that **zsh** has been reinstalled:

```
$ rpm -q zsh
zsh-5.0.2-7.el7.x86_64
```

6. `$ zsh`
....
[coop@q7]/tmp/LFS201%